

**AMITY UNIVERSITY UTTAR
PRADESH
NOIDA**



**AIML (203)
Deep Neural Networks Lab File**

Submitted to:

Prof. M.K. Dutta

Submitted By:

Naman Sharma

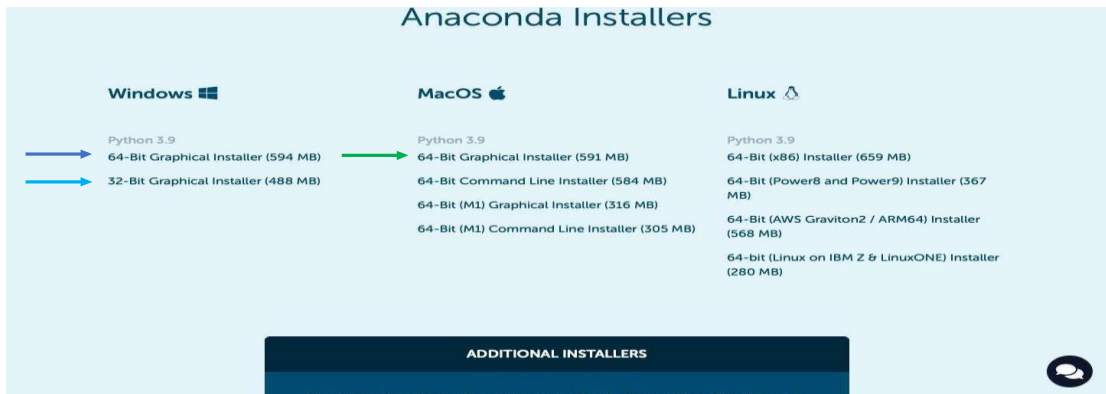
A023119823058

S. No.	Practical	Page no.	Date	Teacher's Signature
1.	Anaconda Installation	1-2	16/12/24	
2.	Bank Churn ANN Algorithm	3-15	23/12/24	
3.	ANN algorithm using bird data	16-37	13/1/25	
4.	Creating CNN model from scratch	38-46	20/1/25	
5.	Deep Learning Training and Architecture	47-69	27/1/25	
6.	Text data handling with RNN for sentiment analysis	70-78	3/2/25	
7.	Sentiment analysis using RNN-LSTM on tweets data	79-91	10/2/25	
8.	Auto encoders using MNIST data	92-97	17/2/25	
9.	Variational Autoencoders	98-106	24/2/25	
10.	Implementing GAN Architecture on MINST dataset	107-116	10/3/25	

Anaconda Installation Guide

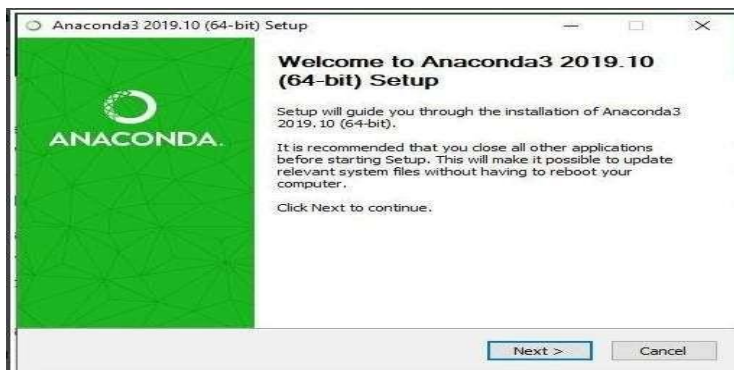
1. Download Anaconda

- Visit the official website: [Anaconda Distribution](https://www.anaconda.com/distribution/).
- Select the appropriate version based on your system (Windows, macOS, Linux).



2. Start Installation

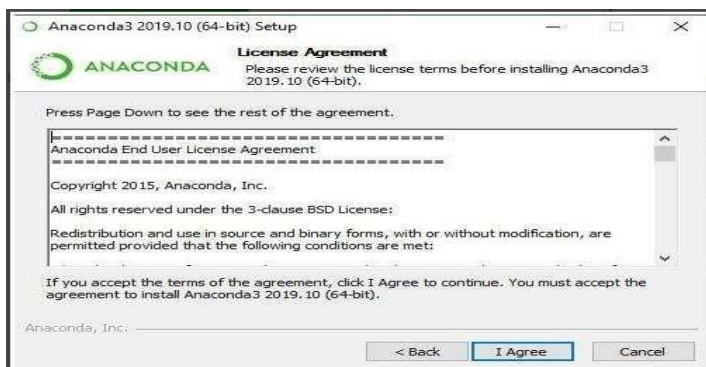
- Run the downloaded installer.



- Click **Next** to begin the installation process.


3. Accept License Agreement

- Read and accept the license agreement to proceed.



4. Select Installation Type

- Choose "**Just Me**" if you want to install for a single user.



5. Choose Installation Location

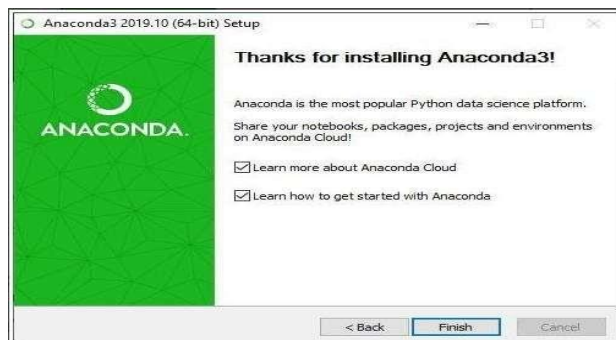
- Select the directory where Anaconda should be installed.

6. Configure Advanced Options (Optional)

- Add Anaconda to your system **PATH** (not recommended).
- Register Anaconda as the default Python environment.

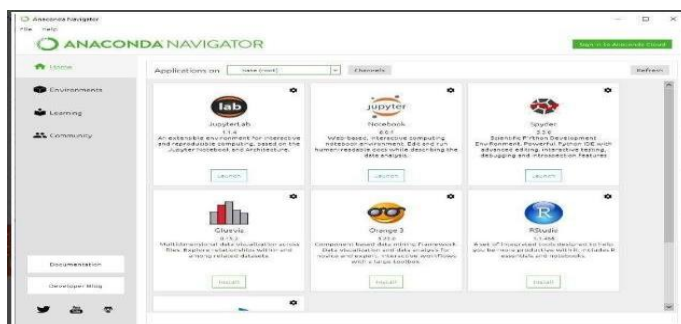
7. Complete Installation

- Click **Install** and wait for the process to finish.
- Click **Finish** once done.



8. Launch Anaconda

- Open **Anaconda Navigator** from the Start Menu.
- You can now start using tools like **Jupyter Notebook** for Python coding.



Experiment - 1

Bank Churn ANN

Importing Necessary Libraries

```
import numpy as np
import pandas as pd
```

Loading the Churn Dataset

```
churn_data = pd.read_csv('/content/Churn_Modelling.csv', delimiter = ',')
churn_data.head(5)
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	\
0	1	15634602	Hargrave	619	France	Female	42	
1	2	15647311	Hill	608	Spain	Female	41	
2	3	15619304	Onio	502	France	Female	42	
3	4	15701354	Boni	699	France	Female	39	
4	5	15737888	Mitchell	850	Spain	Female	43	

	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	\
0	2	0.00	1	1	1	
1	1	83807.86	1	0	1	
2	8	159660.80	3	1	0	
3	1	0.00	2	0	0	
4	2	125510.82	1	1	1	

	EstimatedSalary	Exited
0	101348.88	1
1	112542.58	0
2	113931.57	1
3	93826.63	0
4	79084.10	0

Accessing the Column Names in the Dataset

```
churn_data.columns
```

```
Index(['RowNumber', 'CustomerId', 'Surname', 'CreditScore', 'Geography',  
      'Gender', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'HasCrCard',  
      'IsActiveMember', 'EstimatedSalary', 'Exited'],  
      dtype='object')
```

Setting Column as a Index

```
churn_data = churn_data.set_index('RowNumber')
churn_data.head()
```

\	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure
RowNumber							
1	15634602	Hargrave	619	France	Female	42	2
2	15647311	Hill	608	Spain	Female	41	1
3	15619304	Onio	502	France	Female	42	8
4	15701354	Boni	699	France	Female	39	1
5	15737888	Mitchell	850	Spain	Female	43	2

	Balance	NumOfProducts	HasCrCard	IsActiveMember	\
RowNumber					
1	0.00	1	1	1	
2	83807.86	1	0	1	
3	159660.80	3	1	0	
4	0.00	2	0	0	
5	125510.82	1	1	1	

	EstimatedSalary	Exited
RowNumber		
1	101348.88	1
2	112542.58	0
3	113931.57	1
4	93826.63	0
5	79084.10	0

Finding the Shape of the Dataset

```
churn_data.shape
```

```
(10000, 13)
```

```
churn_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 10000 entries, 1 to 10000
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   CustomerId            10000 non-null  int64
1   Surname                10000 non-null  object
2   CreditScore            10000 non-null  int64
3   Geography              10000 non-null  object
4   Gender                 10000 non-null  object
5   Age                    10000 non-null  int64
6   Tenure                 10000 non-null  int64
7   Balance                10000 non-null  float64
8   NumOfProducts          10000 non-null  int64
9   HasCrCard              10000 non-null  int64
10  IsActiveMember         10000 non-null  int64
11  EstimatedSalary        10000 non-null  float64
12  Exited                  10000 non-null  int64
```

```
dtypes: float64(2), int64(8), object(3)
memory usage: 1.1+ MB
```

Checking Missing Values

```
churn_data.isna().sum()
```

```
CustomerId      0
Surname         0
CreditScore     0
Geography       0
Gender          0
Age            0
Tenure         0
Balance        0
NumOfProducts  0
HasCrCard       0
IsActiveMember  0
EstimatedSalary 0
Exited         0
dtype: int64
```

```
churn_data.nunique()
```

```
CustomerId      10000
Surname         2932
CreditScore     460
Geography       3
Gender          2
Age            70
Tenure         11
Balance        6382
NumOfProducts   4
HasCrCard       2
IsActiveMember  2
EstimatedSalary 9999
Exited         2
dtype: int64
```

```
churn_data.drop(['CustomerId', 'Surname'], axis=1, inplace=True)
```

```
churn_data.head()
```

	CreditScore	Geography	Gender	Age	Tenure	Balance \
RowNumber						
1	619	France	Female	42	2	0.00
2	608	Spain	Female	41	1	83807.86
3	502	France	Female	42	8	159660.80
4	699	France	Female	39	1	0.00
5	850	Spain	Female	43	2	125510.82

	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
RowNumber					
1	1	1		101348.88	1
2	1	0		112542.58	0
3	3	1		113931.57	1
4	2	0		93826.63	0
5	1	1		79084.10	0

```
churn_data.shape
```

```
(10000, 11)
```

Some Visualizations

```
from matplotlib import pyplot as plt
import seaborn as sns
from scipy import stats
df = churn_data.copy()

def plot_univariate(col):
    if(df[col].nunique()>2):
        plt.figure(figsize=(10,7))
        h = 0.15
        rot=90
    else:
        plt.figure(figsize=(6,6))
        h = 0.5
        rot=0
    plot = sns.countplot(x = df[col], palette='pastel')

    for bars in plot.containers:
        for p in bars:
            plot.annotate(format(p.get_height()), (p.get_x() +
p.get_width()*0.5, p.get_height()),
                        ha = 'center', va = 'bottom')
            plot.annotate(f'{p.get_height()*100/df[col].shape[0] : .1f}%',
(p.get_x() + p.get_width()*0.5, h*p.get_height()),
                        ha = 'center', va = 'bottom', rotation=rot)

def spearman(df,hue):
    feature = []
```



```

correlation = []
result = []
for col in df.columns:
    corr, p = stats.spearmanr(df[col], df[hue])
    feature.append(col)
    correlation.append(corr)
    alpha = 0.05
    if p > alpha:
        result.append('No correlation (fail to reject H0)')
    else:
        result.append('Some correlation (reject H0)')
c = pd.DataFrame({'Feature Name':feature,'correlation
coefficient':correlation, 'Inference':result})
display(c)

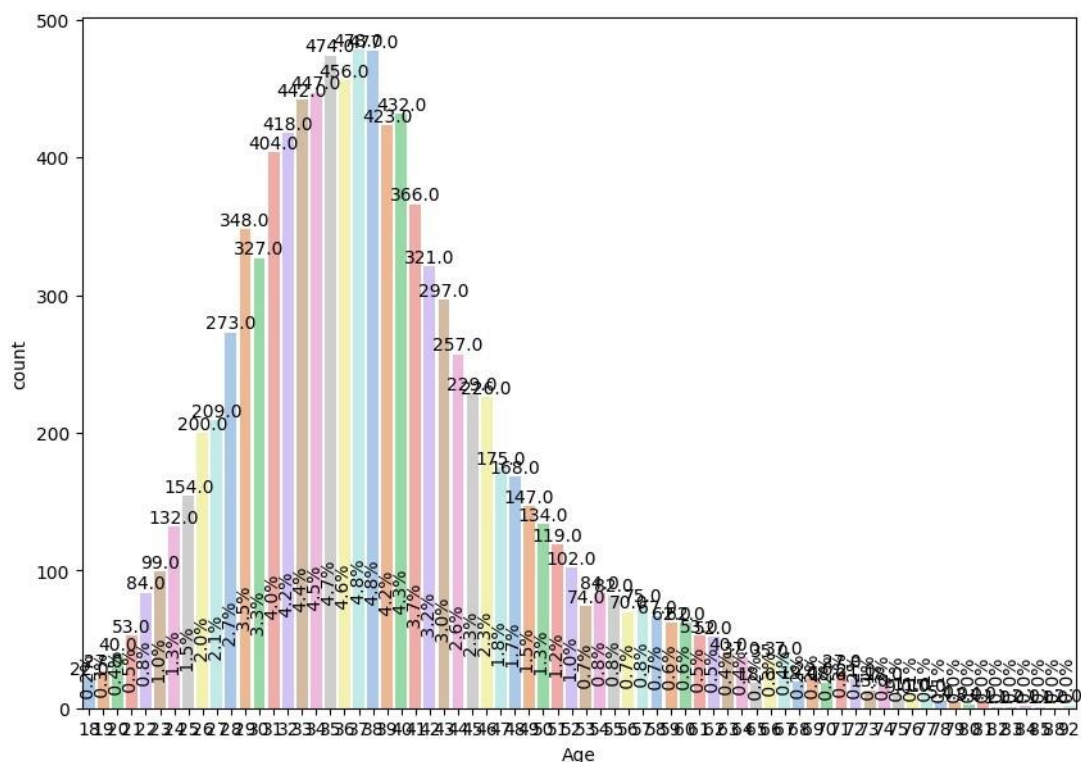
```

```
plot_univariate('Age')
```

/tmp/ipykernel_6944/4200767248.py:10: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
plot = sns.countplot(x = df[col], palette='pastel')
```



```
spearman(churn_data, 'Age')
```

	Feature Name	correlation coefficient \
0	CreditScore	-0.007974
1	Geography	0.035351
2	Gender	-0.029785
3	Age	1.000000
4	Tenure	-0.010405
5	Balance	0.033304
6	NumOfProducts	-0.058566
7	HasCrCard	-0.015278
8	IsActiveMember	0.039839
9	EstimatedSalary	-0.002431
10	Exited	0.323968

	Inference
0	No correlation (fail to reject H0)
1	Some correlation (reject H0)
2	Some correlation (reject H0)
3	Some correlation (reject H0)
4	No correlation (fail to reject H0)
5	Some correlation (reject H0)
6	Some correlation (reject H0)
7	No correlation (fail to reject H0)
8	Some correlation (reject H0)
9	No correlation (fail to reject H0)
10	Some correlation (reject H0)

```

from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
churn_data[['Geography', 'Gender']] = churn_data[['Geography',
'Gender']].apply(le.fit_transform)

churn_data.head()

```

	CreditScore	Geography	Gender	Age	Tenure	Balance \
RowNumber						
1	619	0	0	42	2	0.00
2	608	2	0	41	1	83807.86
3	502	0	0	42	8	159660.80
4	699	0	0	39	1	0.00
5	850	2	0	43	2	125510.82

	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
RowNumber					
1	1	1		101348.88	1
2	1	0		112542.58	0

3	3	1	0	113931.57	1
4	2	0	0	93826.63	0
5	1	1	1	79084.10	0

Seperating Label from Data

```
y = churn_data.Exited
X = churn_data.drop(['Exited'],axis=1)
```

```
X.columns
```

```
Index(['CreditScore', 'Geography', 'Gender', 'Age', 'Tenure', 'Balance',
      'NumOfProducts', 'HasCrCard', 'IsActiveMember', 'EstimatedSalary'],
      dtype='object')
```

```
y
```

```
RowNumber
```

```
1      1
2      0
3      1
4      0
5      0
```

```
..
```

```
9996    0
9997    0
9998    1
9999    1
10000    0
```

```
Name: Exited, Length: 10000, dtype: int64
```

Splitting the Data into Training and Testing

```
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3,
random_state = 2)
```

```
print("Shape of the X_train", X_train.shape)
print("Shape of the X_test", X_test.shape)
print("Shape of the y_train", y_train.shape)
print("Shape of the y_test", y_test.shape)
```

```
Shape of the X_train (7000, 10)
```

```
Shape of the X_test (3000, 10)
```

```
Shape of the y_train (7000,)
```

```
Shape of the y_test (3000,)
```

```
from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)
```

Building the ANN Model

```
from keras.models import Sequential
from keras.layers import Dense
```

```
2024-05-21 08:34:06.981674: E
external/local_xla/xla/stream_executor/cuda/cuda_dnn.cc:9373] Unable to
register cuDNN factory: Attempting to register factory for plugin cuDNN when
one has already been registered
```

```
2024-05-21 08:34:06.981725: E
external/local_xla/xla/stream_executor/cuda/cuda_fft.cc:607] Unable to
register cuFFT factory: Attempting to register factory for plugin cuFFT when
one has already been registered
```

```
2024-05-21 08:34:06.982807: E
external/local_xla/xla/stream_executor/cuda/cuda_blas.cc:1534] Unable to
register cuBLAS factory: Attempting to register factory for plugin cuBLAS
when one has already been registered
```

```
2024-05-21 08:34:06.989509: I
tensorflow/core/platform/cpu_feature_guard.cc:183] This TensorFlow binary is
optimized to use available CPU instructions in performance-critical
operations.
```

To enable the following instructions: SSE3 SSE4.1 SSE4.2 AVX, in other operations, rebuild TensorFlow with the appropriate compiler flags.

```
classifier = Sequential()
# Adding the input layer and the first hidden layer
classifier.add(Dense(units = 8, kernel_initializer = 'uniform', activation =
'relu', input_dim = 10))
```

```
# Adding the second hidden layer
classifier.add(Dense(units = 16, kernel_initializer = 'uniform', activation =
'relu'))
```

```
# Adding the output layer
classifier.add(Dense(units = 1, kernel_initializer = 'uniform', activation =
'sigmoid'))
```

```
2024-05-21 08:34:10.706731: I
tensorflow/core/common_runtime/gpu/gpu_device.cc:1926] Created device
/job:localhost/replica:0/task:0/device:GPU:0 with 17947 MB memory: ->
device: 0, name: NVIDIA A100-SXM4-40GB MIG 3g.20gb, pci bus id: 0000:bd:00.0,
compute capability: 8.0
```

Compiling and Fitting the Model

```
classifier.compile(optimizer = 'adam', loss = 'binary_crossentropy', metrics  
= ['accuracy'])
```

```
# Fitting the ANN to the Training set
```

```
classifier.fit(X_train, y_train, batch_size = 10, epochs = 100, verbose = 1)
```

Epoch 1/100

2024-05-21 08:34:30.051165: I external/local_xla/xla/service/service.cc:168]
XLA service 0x7fc261d536a0 initialized for platform CUDA (this does not
guarantee that XLA will be used). Devices:

2024-05-21 08:34:30.051210: I external/local_xla/xla/service/service.cc:176]
StreamExecutor device (0): NVIDIA A100-SXM4-40GB MIG 3g.20gb, Compute
Capability 8.0

2024-05-21 08:34:30.056612: I
tensorflow/compiler/mlir/tensorflow/utils/dump_mlir_util.cc:269] disabling
MLIR crash reproducer, set env var `MLIR_CRASH_REPRODUCER_DIRECTORY` to
enable.

2024-05-21 08:34:30.108404: I
external/local_xla/xla/stream_executor/cuda/cuda_dnn.cc:467] Loaded cuDNN
version 90000

WARNING: All log messages before absl::InitializeLog() is called are written
to STDERR

I0000 00:00:1716280470.197288 7745 device_compiler.h:186] Compiled cluster
using XLA! This line is logged at most once for the lifetime of the process.

700/700 [=====] - 3s 1ms/step - loss: 0.4773 -
accuracy: 0.7924

Epoch 2/100

700/700 [=====] - 1s 1ms/step - loss: 0.4296 -
accuracy: 0.7926

Epoch 3/100

700/700 [=====] - 1s 1ms/step - loss: 0.4231 -
accuracy: 0.8116

Epoch 4/100

700/700 [=====] - 1s 1ms/step - loss: 0.4188 -
accuracy: 0.8240

Epoch 5/100

700/700 [=====] - 1s 1ms/step - loss: 0.4158 -
accuracy: 0.8297

Epoch 6/100

700/700 [=====] - 1s 1ms/step - loss: 0.4134 -
accuracy: 0.8314

Epoch 7/100

700/700 [=====] - 1s 1ms/step - loss: 0.4120 -
accuracy: 0.8291

```

700/700 [=====] - 1s 1ms/step - loss: 0.3837 -
accuracy: 0.8449
Epoch 92/100
700/700 [=====] - 1s 1ms/step - loss: 0.3820 -
accuracy: 0.8456
Epoch 93/100
700/700 [=====] - 1s 1ms/step - loss: 0.3771 -
accuracy: 0.8496
Epoch 94/100
700/700 [=====] - 1s 1ms/step - loss: 0.3705 -
accuracy: 0.8497
Epoch 95/100
700/700 [=====] - 1s 1ms/step - loss: 0.3610 -
accuracy: 0.8530
Epoch 96/100
700/700 [=====] - 1s 1ms/step - loss: 0.3520 -
accuracy: 0.8581
Epoch 97/100
700/700 [=====] - 1s 1ms/step - loss: 0.3464 -
accuracy: 0.8581
Epoch 98/100
700/700 [=====] - 1s 1ms/step - loss: 0.3428 -
accuracy: 0.8616
Epoch 99/100
700/700 [=====] - 1s 1ms/step - loss: 0.3409 -
accuracy: 0.8613
Epoch 100/100
700/700 [=====] - 1s 1ms/step - loss: 0.3404 -
accuracy: 0.8590

```

<keras.src.callbacks.History at 0x7fcbe2df92d0>

Testing the Model

```

score, acc = classifier.evaluate(X_train, y_train,
                                batch_size=10)
print('Train score:', score)
print('Train accuracy:', acc)

# Predicting the Test set results
y_pred = classifier.predict(X_test)
y_pred = (y_pred > 0.5)

print('*'*20)
score, acc = classifier.evaluate(X_test, y_test,
                                batch_size=10)
print('Test score:', score)
print('Test accuracy:', acc)

700/700 [=====] - 1s 884us/step - loss: 0.3388 -
accuracy: 0.8609

```

```
Train score: 0.33875057101249695
Train accuracy: 0.8608571290969849
94/94 [=====] - 0s 667us/step
*****
300/300 [=====] - 0s 910us/step - loss: 0.3567 -
accuracy: 0.8577
Test score: 0.3566587269306183
Test accuracy: 0.8576666712760925
```

Confusion Matrix


```
from sklearn.metrics import confusion_matrix
target_names = ['Retained', 'Closed']
cm = confusion_matrix(y_test, y_pred)
print(cm)

[[2338   77]
 [ 350  235]]

import matplotlib.pyplot as plt
import seaborn as sns

p = sns.heatmap(pd.DataFrame(cm), annot=True, xticklabels=target_names,
yticklabels=target_names, cmap="YlGnBu", fmt='g')
plt.title('Confusion matrix', y=1.1)
plt.ylabel('Actual label')
plt.xlabel('Predicted label')

Text(0.5, 23.52222222222222, 'Predicted label')
```



Classification

```
#import classification_report
from sklearn.metrics import classification_report
print(classification_report(y_test,y_pred, target_names=target_names))
```

	precision	recall	f1-score	support
Retained	0.87	0.97	0.92	2415
Closed	0.75	0.40	0.52	585
accuracy			0.86	3000
macro avg	0.81	0.68	0.72	3000
weighted avg	0.85	0.86	0.84	3000

#ROC curve

```
from sklearn.metrics import roc_curve, auc
y_pred_proba = classifier.predict(X_test)
fpr, tpr, thresholds = roc_curve(y_test, y_pred_proba)
roc_auc = auc(fpr, tpr)
```




```

plt.plot([0,1],[0,1], 'k--')
plt.plot(fpr, tpr, label='AUC (area = %0.2f)' % roc_auc)
plt.xlabel('fpr')
plt.ylabel('tpr')
plt.grid()
plt.legend(loc="lower right")
plt.title('ROC curve')
plt.show()

```

94/94 [=====] - 0s 667us/step



```

#Area under ROC curve
from sklearn.metrics import roc_auc_score
roc_auc_score(y_test, y_pred_proba)

0.8357169400647662

```

ANN using Bird data

#Importing Necessary Libraries

```
import numpy as np
import pandas as pd
```

#Loading the Birds Dataset

```
bird_data = pd.read_csv('/content/bird.csv', delimiter = ',')
bird_data.head(5)
```

		id				huml	humw	ulnal	ulnaw	feml
femw		tibl	tibw			tarl	tarw	type		
0	0	80.78	6.68	72.01	4.88	41.81	3.70	5.50	4.03	38.70
3.84		SW								
1	1	88.91	6.63	80.53	5.59	47.04	4.30	80.22	4.51	41.50
4.01		SW								
2	2	79.97	6.37	69.26	5.28	43.07	3.90	75.35	4.04	38.31
3.34		SW								
3	3	77.65	5.70	65.76	4.77	40.04	3.52	69.17	3.40	35.78
3.41		SW								
4	4	62.80	4.84	52.09	3.73	33.95	2.72	56.27	2.96	31.88
3.13		SW								

#Accessing the Column Names in the Dataset

```
bird_data.columns
```

```
Index(['id', 'huml', 'humw', 'ulnal', 'ulnaw', 'feml', 'femw', 'tibl',
      'tibw',
      'tarl', 'tarw', 'type'],
      dtype='object')
```

```
bird_data = bird_data.set_index('id')
bird_data.head()
```

[illegible]

```
3    77.65    5.70    65.76    4.77    40.04    3.52    69.17    3.40    35.78    3.41  
SW  
4    62.80    4.84    52.09    3.73    33.95    2.72    56.27    2.96    31.88    3.13  
SW
```

#Finding the Shape of the Dataset

```
bird_data.shape
```

```
(420, 11)
```

```
bird_data.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 420 entries, 0 to 419
Data columns (total 11 columns):
 #   Column  Non-Null Count  Dtype
---  -
 0   huml    419 non-null    float64
 1   humw    419 non-null    float64
 2   ulnal   417 non-null    float64
 3   ulnaw   418 non-null    float64
 4   feml    418 non-null    float64
 5   femw    419 non-null    float64
 6   tibl    418 non-null    float64
 7   tibw    419 non-null    float64
 8   tarl    419 non-null    float64
 9   tarw    419 non-null    float64
10  type    420 non-null    object
dtypes: float64(10), object(1)
memory usage: 39.4+ KB
```

#Checking Missing Values

```
bird_data.isna().sum()

huml      1
humw      1
ulnal     3
ulnaw     2
feml      2
femw      1
tibl      2
tibw      1
tarl      1
tarw      1
type      0
dtype: int64

bird_data.dropna(how='any', inplace=True)

bird_data.isna().sum()

huml      0
humw      0
ulnal     0
ulnaw     0
feml      0
femw      0
tibl      0
tibw      0
```

```
tarl      0
tarw      0
type      0
dtype: int64

bird_data.shape

(413, 11)
```

Unique Values in the Data

```
bird_data.nunique()

huml      403
humw      319
ulnal     394
ulnaw     305
feml      397
femw      287
tibl      401
tibw      283
tarl      403
tarw      277
type      6
dtype: int64

bird_data['type'].unique()

array(['SW', 'W', 'T', 'R', 'P', 'SO'], dtype=object)
```

```
from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
bird_data[['type']] = bird_data[['type']].apply(le.fit_transform)

bird_data.head()
```

					huml	humw	ulnal	ulnaw	feml	femw
	tibl	tibw			tarl	tarw	type			
id										
0	80.78	6.68	72.01	4.88	41.81	3.70	5.50	4.03	38.70	3.84
3										
1	88.91	6.63	80.53	5.59	47.04	4.30	80.22	4.51	41.50	4.01
3										
2	79.97	6.37	69.26	5.28	43.07	3.90	75.35	4.04	38.31	3.34
3										
3	77.65	5.70	65.76	4.77	40.04	3.52	69.17	3.40	35.78	3.41
3										
4	62.80	4.84	52.09	3.73	33.95	2.72	56.27	2.96	31.88	3.13
3										

#Seperating Label from Data

```
y = bird_data['type']
X = bird_data.drop(['type'],axis=1)

X.columns

Index(['huml', 'humw', 'ulnal', 'ulnaw', 'feml', 'femw', 'tibl',
      'tibw',
      'tarl', 'tarw'],
      dtype='object')

y

id
0      3
1      3
2      3
3      3
4      3
..
415    2
416    2
417    2
418    2
419    2
Name: type, Length: 413, dtype: int64

y.shape

(413,)
```

#Splitting the Data into Training and Testing

```
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size =
0.2, random_state = 2)

print("Shape of the X_train", X_train.shape)
print("Shape of the X_test", X_test.shape)
print("Shape of the y_train", y_train.shape)
print("Shape of the y_test", y_test.shape)

Shape of the X_train (330, 10)
Shape of the X_test (83, 10)
Shape of the y_train (330, 6)
Shape of the y_test (83, 6)
```

#Building the ANN Model

sequential model to initialise our ann and dense module to build the layers

```
from keras.models import Sequential
from keras.layers import Dense
```

```
classifier = Sequential()
```

Adding the input layer and the first hidden layer

```
classifier.add(Dense(units = 8, kernel_initializer = 'uniform',
activation = 'relu', input_dim = 10))
```

Adding the second hidden layer

```
classifier.add(Dense(units = 16, kernel_initializer = 'uniform',
activation = 'relu'))
```

Adding the third hidden layer

```
classifier.add(Dense(units = 32, kernel_initializer = 'uniform',
activation = 'relu'))
```

Adding the output layer

```
classifier.add(Dense(units = 6, kernel_initializer = 'uniform',
activation = 'softmax'))
```

```
2024-05-21 08:54:03.283163: I
tensorflow/core/common_runtime/gpu/gpu_device.cc:1926] Created
device /job:localhost/replica:0/task:0/device:GPU:0 with 17947 MB
memory: -> device: 0, name: NVIDIA A100-SXM4-40GB MIG 3g.20gb, pci
bus id: 0000:bd:00.0, compute capability: 8.0
```

Compiling and Fitting the Model

```
classifier.compile(optimizer = 'adam', loss =  
'categorical_crossentropy', metrics = ['accuracy'])  
  
# Fitting the ANN to the Training set  
classifier.fit(X_train, y_train, batch_size = 16, epochs = 800,  
verbose = 1)  
  
Epoch 1/800  
  
2024-05-21 08:54:07.013668: I  
external/local_xla/xla/service/service.cc:168] XLA service  
0x7fd5fd374770 initialized for platform CUDA (this does not guarantee  
that XLA will be used). Devices:  
2024-05-21 08:54:07.013710: I  
external/local_xla/xla/service/service.cc:176] StreamExecutor device  
(0): NVIDIA A100-SXM4-40GB MIG 3g.20gb, Compute Capability 8.0  
2024-05-21 08:54:07.019479: I  
tensorflow/compiler/mlir/tensorflow/utils/dump_mlir_util.cc:269]  
disabling MLIR crash reproducer, set env var
```



```
`MLIR_CRASH_REPRODUCER_DIRECTORY` to enable.
2024-05-21 08:54:07.057513: I
external/local_xla/xla/stream_executor/cuda/cuda_dnn.cc:467] Loaded
cuDNN version 90000
WARNING: All log messages before absl::InitializeLog() is called are
written to STDERR
I0000 00:00:1716281647.144364 13834 device_compiler.h:186] Compiled
cluster using XLA! This line is logged at most once for the lifetime
of the process.

21/21 [=====] - 2s 2ms/step - loss: 1.7886 -
accuracy: 0.2364
Epoch 2/800
21/21 [=====] - 0s 2ms/step - loss: 1.7772 -
accuracy: 0.3364
Epoch 3/800
21/21 [=====] - 0s 2ms/step - loss: 1.7505 -
accuracy: 0.4364
Epoch 4/800
21/21 [=====] - 0s 2ms/step - loss: 1.6777 -
accuracy: 0.4818
Epoch 5/800
21/21 [=====] - 0s 2ms/step - loss: 1.5472 -
accuracy: 0.4909
Epoch 6/800
21/21 [=====] - 0s 2ms/step - loss: 1.4310 -
accuracy: 0.4970
Epoch 7/800
21/21 [=====] - 0s 2ms/step - loss: 1.3724 -
accuracy: 0.5000
Epoch 8/800
21/21 [=====] - 0s 2ms/step - loss: 1.3375 -
accuracy: 0.5030
Epoch 9/800
21/21 [=====] - 0s 2ms/step - loss: 1.3134 -
accuracy: 0.5030
Epoch 10/800
21/21 [=====] - 0s 2ms/step - loss: 1.2927 -
accuracy: 0.5061
Epoch 11/800
21/21 [=====] - 0s 2ms/step - loss: 1.2780 -
accuracy: 0.5091
Epoch 12/800
21/21 [=====] - 0s 2ms/step - loss: 1.2687 -
accuracy: 0.5091
Epoch 13/800
21/21 [=====] - 0s 2ms/step - loss: 1.2571 -
accuracy: 0.5121
Epoch 14/800
21/21 [=====] - 0s 2ms/step - loss: 1.2445 -
```

```
accuracy: 0.5212
Epoch 15/800
21/21 [=====] - 0s 2ms/step - loss: 1.2328 -
accuracy: 0.5242
Epoch 16/800
21/21 [=====] - 0s 2ms/step - loss: 1.2212 -
accuracy: 0.5273
Epoch 17/800
21/21 [=====] - 0s 2ms/step - loss: 1.2172 -
accuracy: 0.5273
Epoch 18/800
21/21 [=====] - 0s 2ms/step - loss: 1.2037 -
accuracy: 0.5303
Epoch 19/800
21/21 [=====] - 0s 2ms/step - loss: 1.1877 -
accuracy: 0.5303
Epoch 20/800
21/21 [=====] - 0s 2ms/step - loss: 1.1763 -
accuracy: 0.5333
Epoch 21/800
21/21 [=====] - 0s 2ms/step - loss: 1.1673 -
accuracy: 0.5364
Epoch 22/800
21/21 [=====] - 0s 2ms/step - loss: 1.1575 -
accuracy: 0.5394
Epoch 23/800
21/21 [=====] - 0s 2ms/step - loss: 1.1437 -
accuracy: 0.5424
Epoch 24/800
21/21 [=====] - 0s 2ms/step - loss: 1.1276 -
accuracy: 0.5424
Epoch 25/800
21/21 [=====] - 0s 2ms/step - loss: 1.1143 -
accuracy: 0.5455
Epoch 26/800
21/21 [=====] - 0s 2ms/step - loss: 1.1014 -
accuracy: 0.5455
Epoch 27/800
21/21 [=====] - 0s 2ms/step - loss: 1.0870 -
accuracy: 0.5455
Epoch 28/800
21/21 [=====] - 0s 2ms/step - loss: 1.0758 -
accuracy: 0.5455
Epoch 29/800
21/21 [=====] - 0s 2ms/step - loss: 1.0680 -
accuracy: 0.5455
Epoch 30/800
21/21 [=====] - 0s 2ms/step - loss: 1.0548 -
accuracy: 0.5485
```

```
accuracy: 0.9606
Epoch 784/800
21/21 [=====] - 0s 2ms/step - loss: 0.1375 -
accuracy: 0.9545
Epoch 785/800
21/21 [=====] - 0s 2ms/step - loss: 0.1342 -
accuracy: 0.9485
Epoch 786/800
21/21 [=====] - 0s 2ms/step - loss: 0.1369 -
accuracy: 0.9545
Epoch 787/800
21/21 [=====] - 0s 2ms/step - loss: 0.1419 -
accuracy: 0.9636
Epoch 788/800
21/21 [=====] - 0s 2ms/step - loss: 0.1409 -
accuracy: 0.9545
Epoch 789/800
21/21 [=====] - 0s 2ms/step - loss: 0.1464 -
accuracy: 0.9485
Epoch 790/800
21/21 [=====] - 0s 2ms/step - loss: 0.1528 -
accuracy: 0.9515
Epoch 791/800
21/21 [=====] - 0s 2ms/step - loss: 0.1446 -
accuracy: 0.9455
Epoch 792/800
21/21 [=====] - 0s 2ms/step - loss: 0.1383 -
accuracy: 0.9576
Epoch 793/800
21/21 [=====] - 0s 2ms/step - loss: 0.1366 -
accuracy: 0.9576
Epoch 794/800
21/21 [=====] - 0s 2ms/step - loss: 0.1394 -
accuracy: 0.9576
Epoch 795/800
21/21 [=====] - 0s 2ms/step - loss: 0.1360 -
accuracy: 0.9545
Epoch 796/800
21/21 [=====] - 0s 2ms/step - loss: 0.1376 -
accuracy: 0.9576
Epoch 797/800
21/21 [=====] - 0s 2ms/step - loss: 0.1394 -
accuracy: 0.9576
Epoch 798/800
21/21 [=====] - 0s 2ms/step - loss: 0.2885 -
accuracy: 0.9212
Epoch 799/800
21/21 [=====] - 0s 2ms/step - loss: 0.2031 -
accuracy: 0.9273
Epoch 800/800
```

```
21/21 [=====] - 0s 2ms/step - loss: 0.1498 - accuracy: 0.9485
```

```
<keras.src.callbacks.History at 0x7fdf783122c0>
```

#Testing the Model

```
score, acc = classifier.evaluate(X_train, y_train,  
                                batch_size=10)
```

```
print('Train score:', score)
```

```
print('Train accuracy:', acc)
```

```
print('*'*20)
```

```
score, acc = classifier.evaluate(X_test, y_test,  
                                batch_size=10)
```

```
print('Test score:', score)
```

```
print('Test accuracy:', acc)
```

```
33/33 [=====] - 0s 999us/step - loss: 0.1421  
- accuracy: 0.9545
```

```
Train score: 0.14205241203308105
```

```
Train accuracy: 0.9545454382896423
```

```
*****
```

```
9/9 [=====] - 0s 1ms/step - loss: 0.4386 -  
accuracy: 0.8916
```

```
Test score: 0.4385785162448883
```

```
Test accuracy: 0.891566276550293
```

```
# Predicting the Test set results
```

```
pred = classifier.predict(X_test)
```

```
print("Y_pred:", pred)
```

```
print("*****")
```

```
y_pred = np.argmax(pred, axis = 1)
```

```
print("Y_pred:", y_pred)
```

```
print("*****")
```

```
print("Y_test:", y_test)
```

```
y_true = np.argmax(y_test, axis = 1)
```

```
print("*****")
```

```
print("Y_test:", y_true)
```

```
3/3 [=====] - 0s 1ms/step
```

```
Y_pred: [[1.13437502e-10 1.58537 41e-13 9.9999762e-01 1.9958 027e-07  
1.19342267e-13 7.11663781e-17]  
[2.60476186e-03 9.92207229e-01 1.64550258e-07 8.11149876e-05  
1.85898723e-04 4.92081419e-03]  
[2.14676633e-02 2.31342129e-02 3.83366924e-03 7.04811055e-06  
9.51290905e-01 2.66507734e-04]  
[5.19312546e-02 6.41205178e-09 9.46038127e-01 2.02322891e-03  
1.39814540e-06 5.98970792e-06]  
[9.30234570e-11 5.98462997e-14 4.77789929e-30 3.15114677e-01  
0.00000000e+00 6.84885323e-01]  
[1.88400315e-06 6.03930058e-21 4.66691310e-15 9.15600002e-01  
3.42472671e-25 8.43981877e-02]  
[1.35823920e-01 5.13227701e-01 2.56901711e-01 3.06145989e-06  
9.39941630e-02 4.94873129e-05]  
[3.26984525e-01 3.14950244e-09 6.71750968e-05 1.57738656e-01  
3.85546102e-08 5.15209615e-01]  
[3.84205058e-02 7.30247200e-01 5.26678741e-06 1.89626201e-06  
1.04787380e-01 1.26537830e-01]  
[3.14090314e-04 8.32000315e-01 4.38185356e-11 1.05659328e-01  
1.07772495e-13 6.20262362e-02]  
[8.62169103e-12 1.65840709e-11 1.00000000e+00 3.86799854e-08  
1.72668073e-13 8.60259400e-18]  
[2.11188064e-11 4.37313232e-20 7.18141360e-29 9.77322459e-01  
0.00000000e+00 2.26775333e-02]  
[4.72837513e-07 3.80029655e-06 9.99995351e-01 3.83343348e-07  
5.30217150e-08 9.50168676e-12]  
[9.80835333e-02 2.61182129e-01 6.15853071e-01 8.95882567e-06  
2.48215068e-02 5.08308731e-05]  
[8.58899879e-15 2.80578638e-17 1.00000000e+00 3.12439283e-08  
6.21814207e-18 7.50081058e-22]  
[1.29032907e-09 3.42193189e-08 1.00000000e+00 5.48537127e-09  
2.02274794e-10 1.47560022e-15]  
[7.71484792e-01 3.33889749e-10 3.85360676e-04 7.93045461e-02  
1.71934516e-06 1.48823529e-01]  
[3.79146299e-38 0.00000000e+00 0.00000000e+00 1.00000000e+00  
0.00000000e+00 2.19223724e-17]  
[1.42191124e-15 1.41321845e-34 3.30129911e-33 9.99977827e-01
```

```

[0. 0. 0. 0. 1. 0.]
[0. 0. 1. 0. 0. 0.]
[0. 0. 0. 0. 0. 1.]
[0. 1. 0. 0. 0. 0.]
[1. 0. 0. 0. 0. 0.]
[0. 0. 1. 0. 0. 0.]
[0. 0. 1. 0. 0. 0.]
[1. 0. 0. 0. 0. 0.]
[0. 0. 0. 0. 0. 1.]
[0. 0. 0. 1. 0. 0.]
[0. 0. 1. 0. 0. 0.]
[0. 1. 0. 0. 0. 0.]
[0. 0. 0. 0. 0. 1.]
[0. 0. 0. 1. 0. 0.]
[0. 1. 0. 0. 0. 0.]
[0. 0. 0. 1. 0. 0.]
[0. 0. 0. 0. 0. 1.]
[0. 0. 0. 0. 0. 1.]
[0. 0. 1. 0. 0. 0.]
[0. 0. 0. 1. 0. 0.]
[0. 0. 1. 0. 0. 0.]
[0. 0. 1. 0. 0. 0.]
[0. 0. 1. 0. 0. 0.]]
*****
Y_test: [2 1 4 0 5 3 1 5 1 1 2 3 2 0 2 2 5 3 3 4 2 3 2 3 2 3 2 2 5 1
2 3 5 5 3 3
2 2 2 0 3 3 1 2 3 3 1 2 0 3 2 3 1 5 2 2 2 2 2 4 2 5 1 0 2 2 0 5 3 2 1
5 3
1 3 5 5 2 3 2 2 2]

# Making the Confusion Matrix
from sklearn.metrics import confusion_matrix
cm = confusion_matrix(y_true, y_pred)
target_names = ['P', 'R', 'SO', 'SW', 'T', 'W']

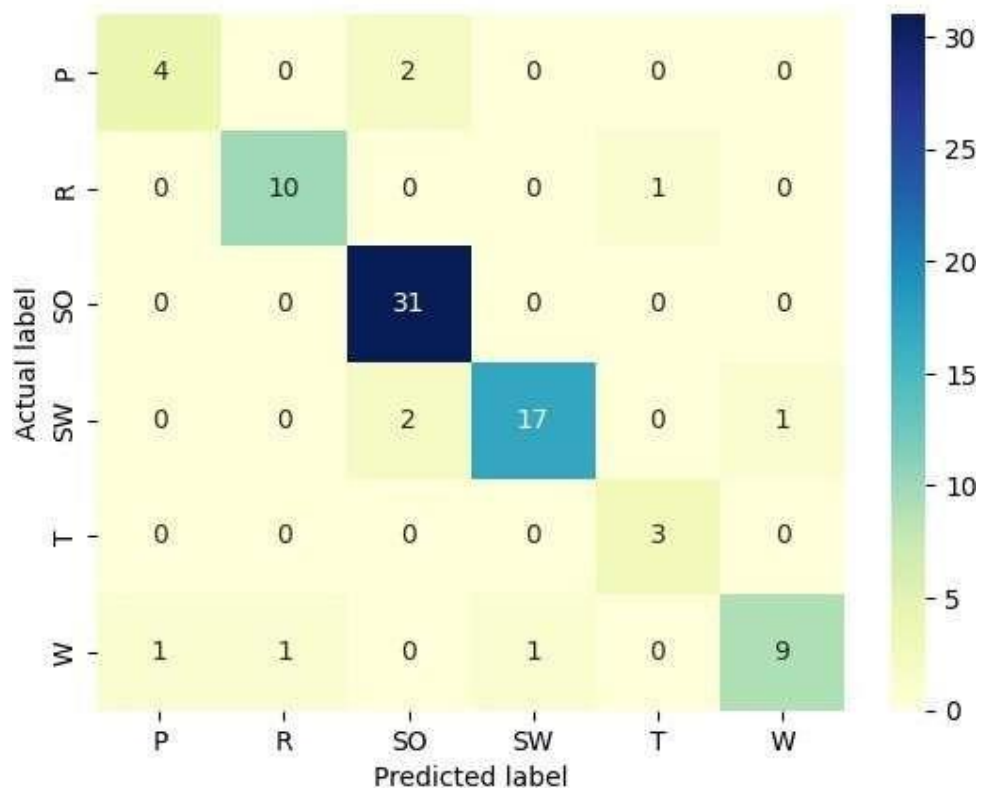
import matplotlib.pyplot as plt
import seaborn as sns

p = sns.heatmap(pd.DataFrame(cm), annot=True, xticklabels=target_names,
yticklabels=target_names, cmap="YlGnBu", fmt='g')
plt.title('Confusion matrix', y=1.1)
plt.ylabel('Actual label')
plt.xlabel('Predicted label')

Text(0.5, 23.52222222222222, 'Predicted label')

```

Confusion matrix



```
#import classification_report
from sklearn.metrics import classification_report
print(classification_report(y_true,y_pred, target_names =
target_names))
```

	precision	recall	f1-score	support
P	0.80	0.67	0.73	6
R	0.91	0.91	0.91	11
SO	0.89	1.00	0.94	31
SW	0.94	0.85	0.89	20
T	0.75	1.00	0.86	3
W	0.90	0.75	0.82	12
accuracy			0.89	83
macro avg	0.86	0.86	0.86	83
weighted avg	0.89	0.89	0.89	83

#

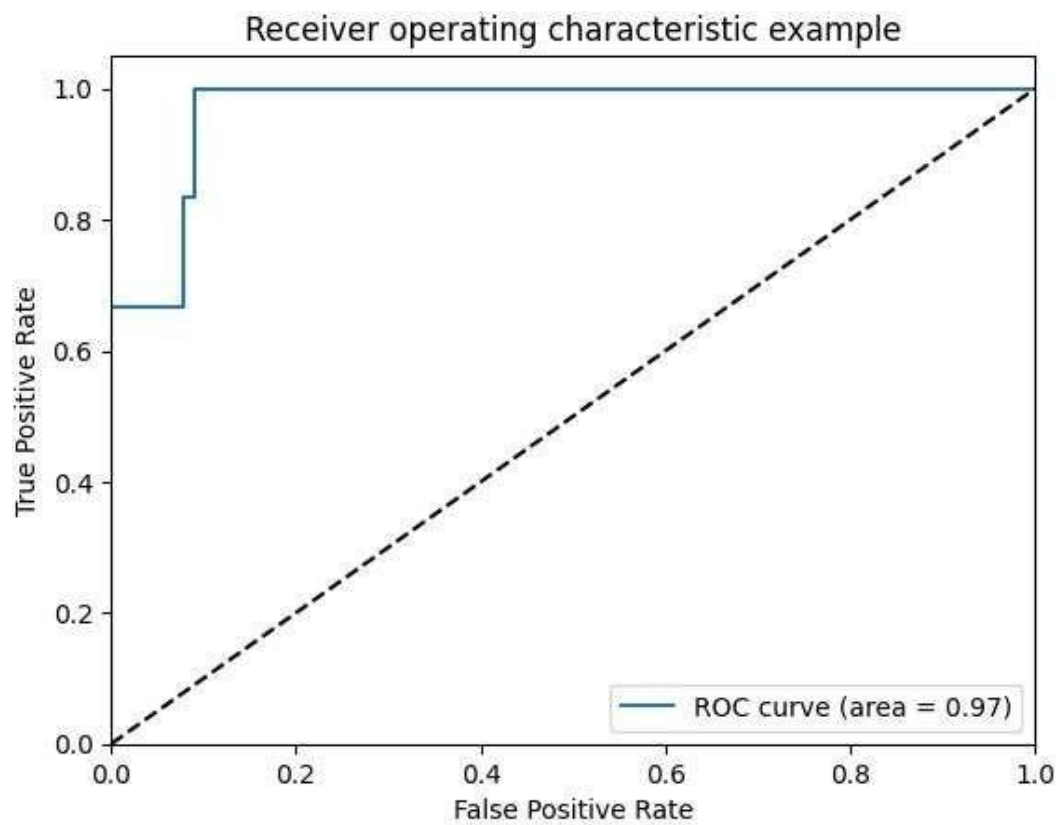
ROC curve

```
from sklearn.metrics import roc_curve, auc
from itertools import cycle
fpr = dict()
tpr = dict()
roc_auc = dict()
for i in range(6):
    fpr[i], tpr[i], _ = roc_curve(y_test[:, i], pred[:, i])
    roc_auc[i] = auc(fpr[i], tpr[i])
```




```
# Plot of a ROC curve for a specific class
```


```
for i in range(6):  
    plt.figure()  
    plt.plot(fpr[i], tpr[i], label='ROC curve (area = %0.2f)' %  
roc_auc[i])  
    plt.plot([0, 1], [0, 1], 'k--')  
    plt.xlim([0.0, 1.0])  
    plt.ylim([0.0, 1.05])  
    plt.xlabel('False Positive Rate')  
    plt.ylabel('True Positive Rate')  
    plt.title('Receiver operating characteristic example')  
    plt.legend(loc="lower right")  
    plt.show()
```

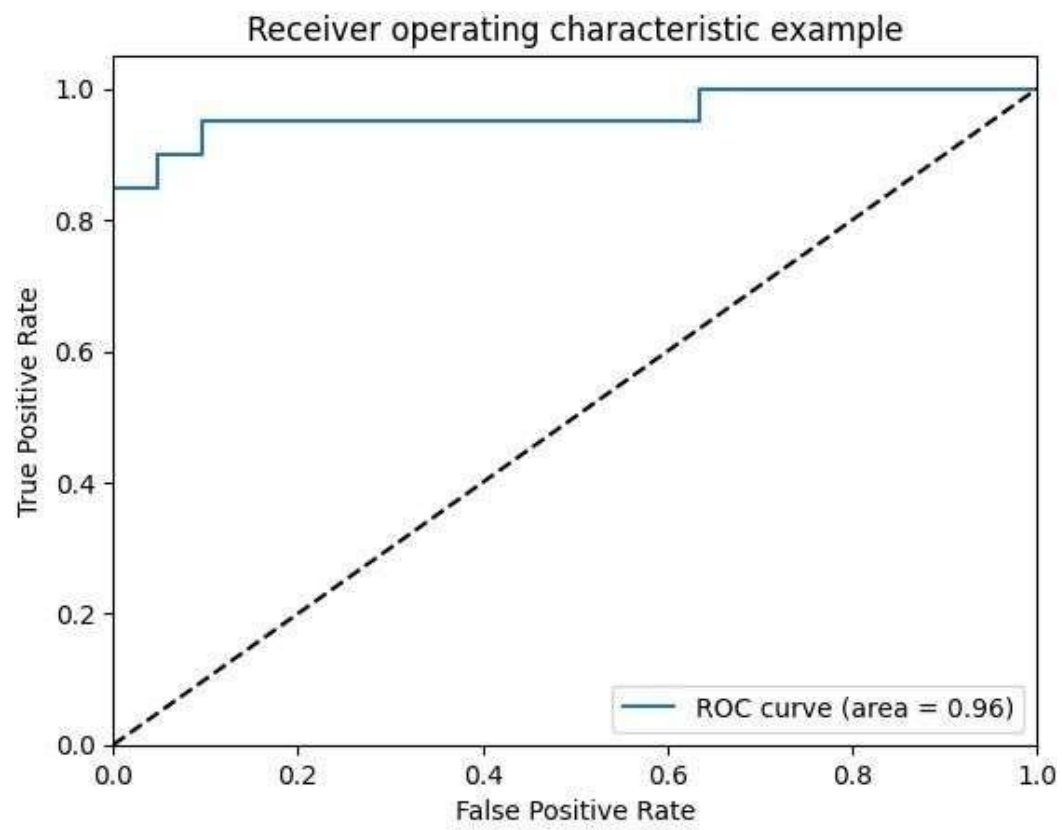


Receiver operating characteristic example




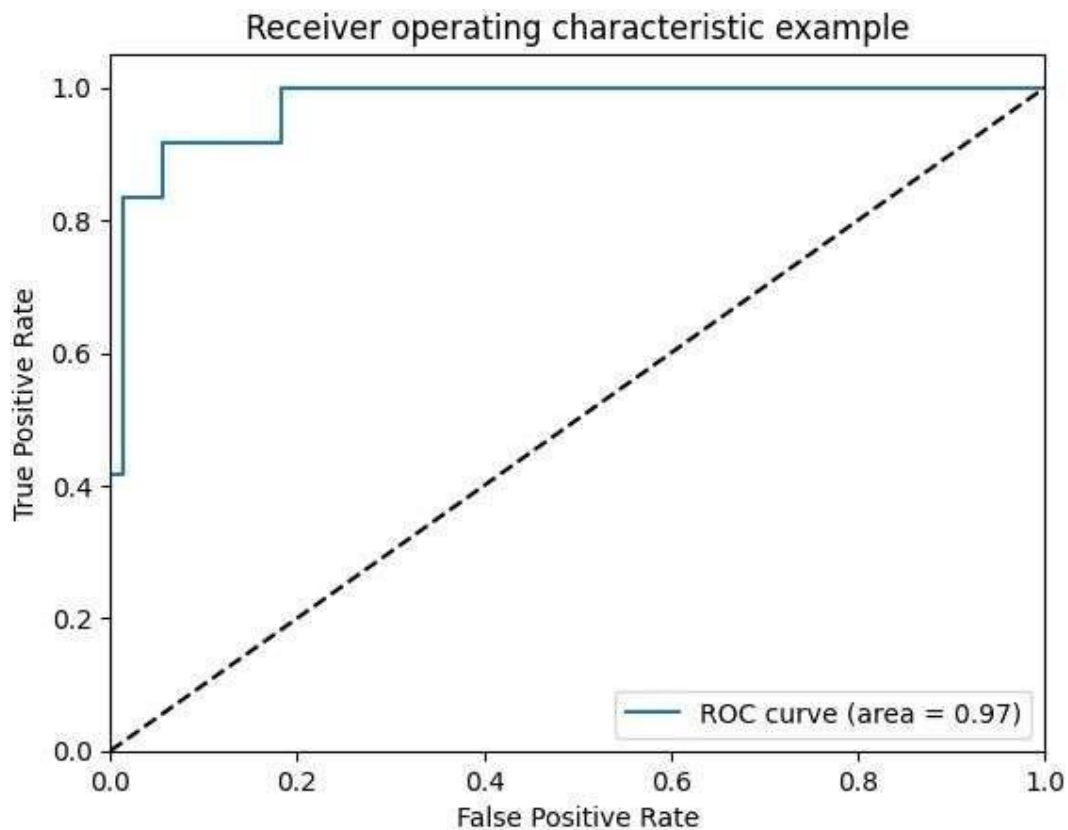
Receiver operating characteristic example






Receiver operating characteristic example





```
fpr = dict()
tpr = dict()
roc_auc = dict()
lw=2
for i in range(6):
    fpr[i], tpr[i], _ = roc_curve(y_test[:, i], pred[:, i])
    roc_auc[i] = auc(fpr[i], tpr[i])
colors = cycle(['blue', 'green', 'red', 'darkorange', 'olive', 'purple'])
for i, color in zip(range(6), colors):
    plt.plot(fpr[i], tpr[i], color=color, lw=2,
             label='AUC = {1:0.4f}'
             ''.format(i, roc_auc[i]))
plt.plot([0, 1], [0, 1], 'k--', lw=lw)
plt.xlim([-0.05, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate', fontsize=15)
plt.ylabel('True Positive Rate', fontsize=15)
# plt.title('Receiver operating characteristic for multi-class data')
plt.legend(loc="lower right")
plt.show()
```



Experiment - 3

Creating CNN Models from Scratch, Compiling of CNN Models, Training and Testing, Plotting of Curves

Importing Necessary libraries

```
from tensorflow.keras import applications
from tensorflow.keras import optimizers
from tensorflow.keras.models import Sequential, Model
from tensorflow.keras.layers import Dropout, Flatten, Dense,
GlobalAveragePooling2D
from tensorflow.keras import backend as k
from tensorflow.keras.callbacks import ModelCheckpoint,
LearningRateScheduler, TensorBoard, EarlyStopping

import numpy as np
from tensorflow.keras import models
import matplotlib.pyplot as plt
from tensorflow.keras.preprocessing import image
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.models import Model
from tensorflow.keras import layers
from tensorflow.keras.layers import Dense, Flatten, Conv2D, MaxPooling2D
from tensorflow.keras import Input
```

Loading the training and testing data and defining the basic parameters

```
train_datagen = ImageDataGenerator(rescale=1./255) # vertical_flip=True,
                                                    # horizontal_flip=True,
                                                    # height_shift_range=0.1,
                                                    # width_shift_range=0.1

validation_datagen = ImageDataGenerator(rescale=1./255)
test_datagen = ImageDataGenerator(rescale=1./255)

# Read the training sample and set the batch size
train_generator = train_datagen.flow_from_directory(
    'C://Users//abhia//Downloads//plant_village
(1)//plant_village/t/rain',
    target_size=(64, 64),
    batch_size=16,
    class_mode='categorical')

# Read Validation data from directory and define target size with batch size
validation_generator = validation_datagen.flow_from_directory(
    '/workspace/Bootcamp/Data/plant_village/val/',
    target_size=(64, 64),
    batch_size=16,
```




```
class_mode='categorical',
shuffle=False)

test_generator = test_datagen.flow_from_directory(
    '/workspace/Bootcamp/Data/plant_village/test/',
    target_size=(64, 64),
    batch_size=1,
    class_mode='categorical',
    shuffle=False)
```

Visualization of few images

```
plt.figure(figsize=(16, 16))
for i in range(1, 17):
    plt.subplot(4, 4, i)
    img, label = test_generator.next()
    # print(img.shape)
    # print(label)
    plt.imshow(img[0])
plt.show()
```



```
img, label = test_generator.next()
img[0].shape
```

```
(64, 64, 3)
```

```
# Create the model
```

```
model = models.Sequential()
```

```
# Add new Layers
```

```
model.add(Conv2D(128, kernel_size=(3,3), activation = 'relu',
input_shape=(64,64,3)))
```

```
model.add(MaxPooling2D(pool_size=(2,2)))
```

```
model.add(Conv2D(64, kernel_size=(3,3), activation = 'relu'))
```

```
model.add(MaxPooling2D(pool_size=(2,2)))
```

```
model.add(Conv2D(64, kernel_size=(3,3), activation = 'relu'))
```

```
model.add(MaxPooling2D(pool_size=(2,2)))
```

```
model.add(layers.Flatten())
```

```
model.add(layers.Dense(32, activation='relu'))
```

```
model.add(layers.Dense(4, activation='softmax'))
model.summary()
```

Compiling and Training the Model

```
model.compile(optimizer = optimizers.Adam(learning_rate = 0.0001),
loss='categorical_crossentropy', metrics=['acc'])
# Train the model
history = model.fit(train_generator,
                    steps_per_epoch=train_generator.samples/train_generator.batch_size,
                    epochs=30,
                    validation_data=validation_generator,

validation_steps=validation_generator.samples/validation_generator.batch_size
,
                    verbose=2)
```

```
Epoch 1/30
Epoch 1/30
188/187 - 2s - loss: 0.0354 - acc: 0.9893 - val_loss: 0.2115 - val_acc:
0.9402
Epoch 2/30
Epoch 1/30
188/187 - 2s - loss: 0.0313 - acc: 0.9890 - val_loss: 0.2183 - val_acc:
0.9339
Epoch 3/30
Epoch 1/30
188/187 - 2s - loss: 0.0292 - acc: 0.9903 - val_loss: 0.2038 - val_acc:
0.9465
Epoch 4/30
Epoch 1/30
188/187 - 2s - loss: 0.0314 - acc: 0.9907 - val_loss: 0.2843 - val_acc:
0.9181
Epoch 5/30
Epoch 1/30
188/187 - 2s - loss: 0.0213 - acc: 0.9940 - val_loss: 0.3160 - val_acc:
0.9181
Epoch 6/30
Epoch 1/30
188/187 - 2s - loss: 0.0276 - acc: 0.9923 - val_loss: 0.2019 - val_acc:
0.9449
Epoch 7/30
Epoch 1/30
188/187 - 2s - loss: 0.0184 - acc: 0.9960 - val_loss: 0.2326 - val_acc:
0.9307
Epoch 8/30
Epoch 1/30
188/187 - 2s - loss: 0.0161 - acc: 0.9970 - val_loss: 0.2079 - val_acc:
0.9386
Epoch 9/30
```

Epoch 1/30
188/187 - 2s - loss: 0.0145 - acc: 0.9980 - val_loss: 0.2025 - val_acc:
0.9465
Epoch 10/30
Epoch 1/30
188/187 - 2s - loss: 0.0180 - acc: 0.9953 - val_loss: 0.2236 - val_acc:
0.9386
Epoch 11/30
Epoch 1/30
188/187 - 2s - loss: 0.0152 - acc: 0.9970 - val_loss: 0.2413 - val_acc:
0.9339
Epoch 12/30
Epoch 1/30
188/187 - 2s - loss: 0.0245 - acc: 0.9913 - val_loss: 0.1965 - val_acc:
0.9465
Epoch 13/30
Epoch 1/30
188/187 - 2s - loss: 0.0202 - acc: 0.9940 - val_loss: 0.3189 - val_acc:
0.9228
Epoch 14/30
Epoch 1/30
188/187 - 2s - loss: 0.0136 - acc: 0.9970 - val_loss: 0.1991 - val_acc:
0.9449
Epoch 15/30
Epoch 1/30
188/187 - 2s - loss: 0.0083 - acc: 0.9983 - val_loss: 0.2098 - val_acc:
0.9402
Epoch 16/30
Epoch 1/30
188/187 - 2s - loss: 0.0081 - acc: 0.9993 - val_loss: 0.2170 - val_acc:
0.9496
Epoch 17/30
Epoch 1/30
188/187 - 2s - loss: 0.0104 - acc: 0.9980 - val_loss: 0.2084 - val_acc:
0.9480
Epoch 18/30
Epoch 1/30
188/187 - 2s - loss: 0.0070 - acc: 0.9993 - val_loss: 0.1953 - val_acc:
0.9480
Epoch 19/30
Epoch 1/30
188/187 - 2s - loss: 0.0578 - acc: 0.9804 - val_loss: 0.4641 - val_acc:
0.8898
Epoch 20/30
Epoch 1/30
188/187 - 2s - loss: 0.0172 - acc: 0.9950 - val_loss: 0.2292 - val_acc:
0.9417
Epoch 21/30
Epoch 1/30
188/187 - 2s - loss: 0.0063 - acc: 0.9993 - val_loss: 0.2090 - val_acc:

```
0.9449
Epoch 22/30
Epoch 1/30
188/187 - 2s - loss: 0.0047 - acc: 1.0000 - val_loss: 0.1988 - val_acc:
0.9449
Epoch 23/30
Epoch 1/30
188/187 - 2s - loss: 0.0035 - acc: 1.0000 - val_loss: 0.2330 - val_acc:
0.9449
Epoch 24/30
Epoch 1/30
188/187 - 2s - loss: 0.0031 - acc: 0.9997 - val_loss: 0.2031 - val_acc:
0.9496
Epoch 25/30
Epoch 1/30
188/187 - 2s - loss: 0.0037 - acc: 1.0000 - val_loss: 0.2201 - val_acc:
0.9480
Epoch 26/30
Epoch 1/30
188/187 - 2s - loss: 0.0041 - acc: 0.9990 - val_loss: 0.2835 - val_acc:
0.9276
Epoch 27/30
Epoch 1/30
188/187 - 2s - loss: 0.0510 - acc: 0.9827 - val_loss: 0.2389 - val_acc:
0.9386
Epoch 28/30
Epoch 1/30
188/187 - 2s - loss: 0.0063 - acc: 0.9993 - val_loss: 0.2344 - val_acc:
0.9465
Epoch 29/30
Epoch 1/30
188/187 - 2s - loss: 0.0030 - acc: 1.0000 - val_loss: 0.2219 - val_acc:
0.9480
Epoch 30/30
Epoch 1/30
188/187 - 2s - loss: 0.0029 - acc: 1.0000 - val_loss: 0.2166 - val_acc:
0.9449
```

```
model.save("CONV_plant_deseas.h5")
print("Saved model to disk")
```

Saved model to disk

```
model = models.load_model('CONV_plant_deseas.h5')
```

Visualization of Accuracy and Loss Curves

```
train_acc = history.history['acc']
val_acc = history.history['val_acc']
train_loss = history.history['loss']
val_loss = history.history['val_loss']

epochs = range(len(train_acc))
plt.plot(epochs, train_acc, 'b', label='Training Accuracy')
plt.plot(epochs, val_acc, 'g', label='Validation Accuracy')
plt.title('Training and Validation Accuracy')
plt.grid()
plt.legend()
plt.figure()
plt.show()

plt.plot(epochs, train_loss, 'b', label='Training Loss')
plt.plot(epochs, val_loss, 'g', label='Validation Loss')
plt.title('Training and Validation Loss')
plt.grid()
plt.legend()
plt.show()
```

Prediction

```
fnames = test_generator_filenames
ground_truth = test_generator_classes
label2index = test_generator_class_indices

idx2label = dict((v,k) for k,v in label2index.items())

# Get the predictions from the model using the generator
predictions = model.predict_generator(test_generator,
steps=test_generator.samples/test_generator.batch_size,verbose=1)
predicted_classes = np.argmax(predictions,axis=1)

errors = np.where(predicted_classes != ground_truth)[0]
print("No of errors = {}/{}".format(len(errors),test_generator.samples))

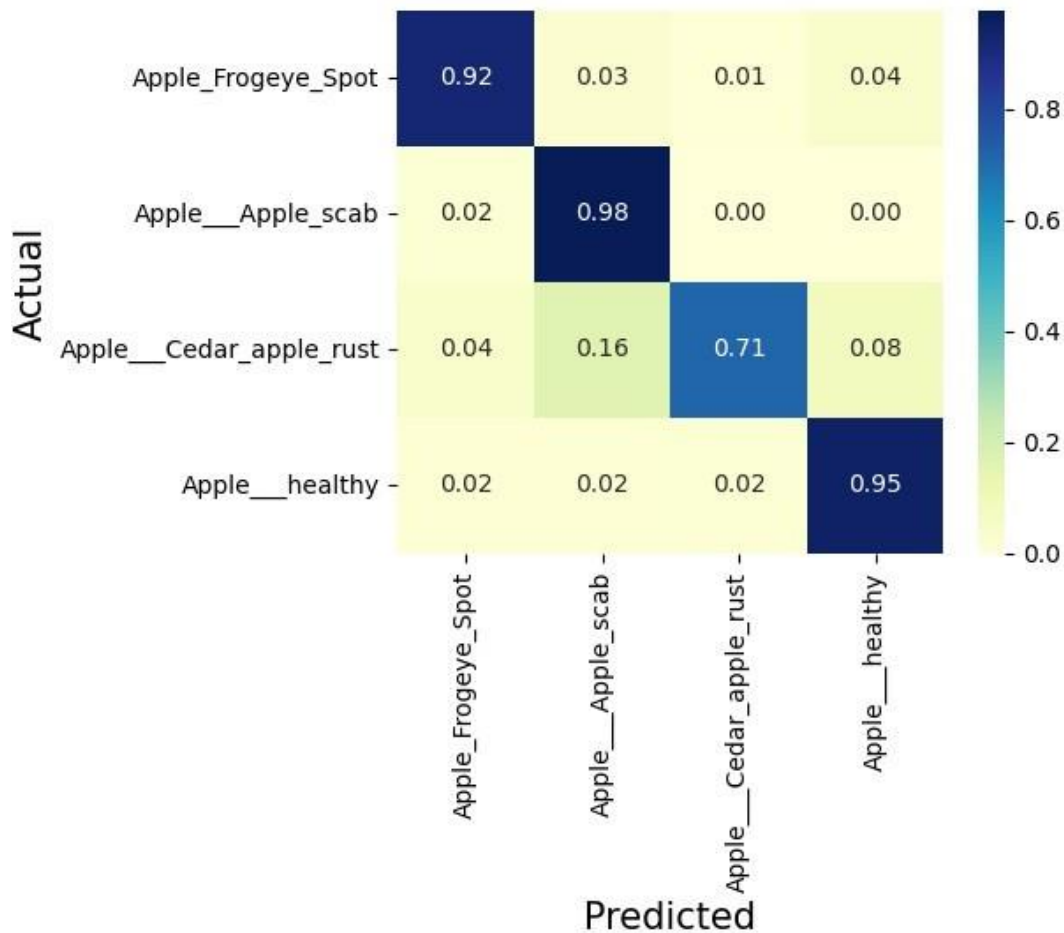
546/546 [=====] - 1s 1ms/step
No of errors = 38/546

accuracy = ((test_generator.samples-len(errors))/test_generator.samples) *
100
accuracy

93.04029304029304
```

Confusion Matrix

```
from sklearn.metrics import confusion_matrix
import seaborn as sns
import numpy as np
from matplotlib import pyplot as plt
cm = confusion_matrix(y_true=ground_truth, y_pred=predicted_classes)
cm = np.array(cm)
# Normalise
cmn = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
fig, ax = plt.subplots(figsize=(5,4))
sns.heatmap(cmn, annot=True, fmt='.2f', xticklabels=label2index,
yticklabels=label2index, cmap="YlGnBu")
plt.ylabel('Actual', fontsize=15)
plt.xlabel('Predicted', fontsize=15)
plt.show(block=False)
```



Classification Report

```
from sklearn.metrics import classification_report
print(classification_report(ground_truth, predicted_classes,
target_names=label2index))
```

	precision	recall	f1-score	support
Apple_Frogeye_Spot	0.91	0.92	0.92	103
Apple___Apple_scab	0.90	0.98	0.94	134
Apple___Cedar_apple_rust	0.85	0.71	0.78	49
Apple___healthy	0.97	0.95	0.96	260
accuracy			0.93	546
macro avg	0.91	0.89	0.90	546
weighted avg	0.93	0.93	0.93	546

Experiment - 4

Deep Learning Training and Architecture, Feature Extraction, Models training with some pretrained models.

Importing Necessary libraries

```
import numpy as np
from tensorflow.keras import Input
from tensorflow.keras import models
from tensorflow.keras import layers
from tensorflow.keras import optimizers
from tensorflow.keras.models import Model
from tensorflow.keras import applications
from tensorflow.keras import backend as k
import matplotlib.pyplot as plt
from tensorflow.keras.optimizers import SGD, Adam
from tensorflow.keras.layers import Dense, Flatten, Conv2D, MaxPooling2D
from tensorflow.keras.preprocessing import image
from tensorflow.keras.models import Sequential, Model
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.layers import Dropout, Flatten, Dense,
GlobalAveragePooling2D
from tensorflow.keras.callbacks import ModelCheckpoint,
LearningRateScheduler, TensorBoard, EarlyStopping
```

Loading the Training and Testing Data and Defining the Basic Parameters

```
train_datagen = ImageDataGenerator(rescale=1./255) # vertical_flip=True,
                                                    # horizontal_flip=True,
                                                    # height_shift_range=0.1,
                                                    # width_shift_range=0.1

validation_datagen = ImageDataGenerator(rescale=1./255)
test_datagen = ImageDataGenerator(rescale=1./255)

# Read the training sample and set the batch size
train_generator = train_datagen.flow_from_directory(
    '/workspace/Bootcamp/Data/plant_village/train/',
    target_size=(128, 128),
    batch_size=16,
    class_mode='categorical')

# Read Validation data from directory and define target size with batch size
validation_generator = validation_datagen.flow_from_directory(
    '/content/plant_village/val/')
```


```
target_size=(128, 128),
batch_size=16,
class_mode='categorical',
shuffle=False)

test_generator = test_datagen.flow_from_directory(
    '/content/plant_village/test/',
    target_size=(128, 128),
    batch_size=1,
    class_mode='categorical',
    shuffle=False)
```

Found 3004 images belonging to 4 classes.
Found 635 images belonging to 4 classes.
Found 547 images belonging to 4 classes.

Visualization of Few Images

```
plt.figure(figsize=(16, 16))
for i in range(1, 17):
    plt.subplot(4, 4, i)
    img, label = test_generator.next()
    # print(img.shape)
    # print(label)
    plt.imshow(img[0])
plt.show()
```



```
img, label = test_generator.next()
img[0].shape
```

```
(128, 128, 3)
```

Exploring Keras Applications for Transfer Learning

```
from tensorflow.keras.applications.vgg16 import VGG16
```

```
## Loading VGG16 model
```

```
base_model = VGG16(weights="imagenet", include_top=False, input_shape= (128, 128, 3))
```

```
# Include_top = False means excluding the model fully connected layers
```

```
base_model.trainable = False ## Not trainable weights,
```

#weights of the VGG16 model will not be updated during training
base_model.summary()

Model: "vgg16"

Layer (type)	Output Shape	Param #
=====		
input_2 (InputLayer)	[(None, 128, 128, 3)]	0
block1_conv1 (Conv2D)	(None, 128, 128, 64)	1792
block1_conv2 (Conv2D)	(None, 128, 128, 64)	36928
block1_pool (MaxPooling2D)	(None, 64, 64, 64)	0
block2_conv1 (Conv2D)	(None, 64, 64, 128)	73856
block2_conv2 (Conv2D)	(None, 64, 64, 128)	147584
block2_pool (MaxPooling2D)	(None, 32, 32, 128)	0
block3_conv1 (Conv2D)	(None, 32, 32, 256)	295168
block3_conv2 (Conv2D)	(None, 32, 32, 256)	590080
block3_conv3 (Conv2D)	(None, 32, 32, 256)	590080
block3_pool (MaxPooling2D)	(None, 16, 16, 256)	0
block4_conv1 (Conv2D)	(None, 16, 16, 512)	1180160
block4_conv2 (Conv2D)	(None, 16, 16, 512)	2359808
block4_conv3 (Conv2D)	(None, 16, 16, 512)	2359808
block4_pool (MaxPooling2D)	(None, 8, 8, 512)	0
block5_conv1 (Conv2D)	(None, 8, 8, 512)	2359808
block5_conv2 (Conv2D)	(None, 8, 8, 512)	2359808
block5_conv3 (Conv2D)	(None, 8, 8, 512)	2359808
block5_pool (MaxPooling2D)	(None, 4, 4, 512)	0
=====		
Total params: 14,714,688		
Trainable params: 0		
Non-trainable params: 14,714,688		

Adding top layers according to number of classes in our data

```
flatten_layer = layers.GlobalAveragePooling2D()
# dense_layer_1 = layers.Dense(64, activation='relu')
# dense_layer_2 = layers.Dense(32, activation='relu')
prediction_layer = layers.Dense(4, activation='softmax')
```

```
model = models.Sequential([
    base_model,
    flatten_layer,
    prediction_layer
])
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
vgg16 (Model)	(None, 4, 4, 512)	14714688

global_average_pooling2d (Gl	(None, 512)	0

dense (Dense)	(None, 4)	2052
=====		
Total params: 14,716,740		
Trainable params: 2,052		
Non-trainable params: 14,714,688		

Training

```
# sgd = SGD(lr=0.001,decay=1e-6, momentum=0.9, nesterov=True)
# We are going to use accuracy metrics and cross entropy loss as performance
# parameters
model.compile(optimizer = Adam(learning_rate = 0.0001),
loss='categorical_crossentropy', metrics=['acc'])
# Train the model
history = model.fit(train_generator,
    steps_per_epoch=train_generator.samples/train_generator.batch_size,
    epochs=30,
    validation_data=validation_generator,

validation_steps=validation_generator.samples/validation_generator.batch_size
,
    verbose=1)

Epoch 1/30
186/187 [=====>.] - ETA: 0s - loss: 0.1192 - acc:
0.9707Epoch 1/30
188/187 [=====] - 6s 34ms/step - loss: 0.1190 - acc:
```

0.9707 - val_loss: 0.1386 - val_acc: 0.9559
Epoch 2/30
186/187 [=====>.] - ETA: 0s - loss: 0.1190 - acc:
0.9701Epoch 1/30
188/187 [=====] - 6s 32ms/step - loss: 0.1184 - acc:
0.9704 - val_loss: 0.1364 - val_acc: 0.9559
Epoch 3/30
186/187 [=====>.] - ETA: 0s - loss: 0.1184 - acc:
0.9711Epoch 1/30
188/187 [=====] - 6s 32ms/step - loss: 0.1180 - acc:
0.9714 - val_loss: 0.1360 - val_acc: 0.9559
Epoch 4/30
186/187 [=====>.] - ETA: 0s - loss: 0.1170 - acc:
0.9704Epoch 1/30
188/187 [=====] - 6s 32ms/step - loss: 0.1175 - acc:
0.9704 - val_loss: 0.1361 - val_acc: 0.9559
Epoch 5/30
186/187 [=====>.] - ETA: 0s - loss: 0.1168 - acc:
0.9694Epoch 1/30
188/187 [=====] - 6s 32ms/step - loss: 0.1172 - acc:
0.9690 - val_loss: 0.1367 - val_acc: 0.9575
Epoch 6/30
186/187 [=====>.] - ETA: 0s - loss: 0.1169 - acc:
0.9704Epoch 1/30
188/187 [=====] - 6s 32ms/step - loss: 0.1169 - acc:
0.9704 - val_loss: 0.1355 - val_acc: 0.9575
Epoch 7/30
186/187 [=====>.] - ETA: 0s - loss: 0.1162 - acc:
0.9704Epoch 1/30
188/187 [=====] - 6s 32ms/step - loss: 0.1164 - acc:
0.9704 - val_loss: 0.1348 - val_acc: 0.9559
Epoch 8/30
186/187 [=====>.] - ETA: 0s - loss: 0.1162 - acc:
0.9707Epoch 1/30
188/187 [=====] - 6s 32ms/step - loss: 0.1164 - acc:
0.9707 - val_loss: 0.1354 - val_acc: 0.9559
Epoch 9/30
186/187 [=====>.] - ETA: 0s - loss: 0.1163 - acc:
0.9711Epoch 1/30
188/187 [=====] - 6s 32ms/step - loss: 0.1162 - acc:
0.9710 - val_loss: 0.1349 - val_acc: 0.9575
Epoch 10/30
186/187 [=====>.] - ETA: 0s - loss: 0.1155 - acc:
0.9704Epoch 1/30
188/187 [=====] - 6s 32ms/step - loss: 0.1157 - acc:
0.9704 - val_loss: 0.1349 - val_acc: 0.9575
Epoch 11/30
186/187 [=====>.] - ETA: 0s - loss: 0.1157 - acc:
0.9711Epoch 1/30
188/187 [=====] - 6s 32ms/step - loss: 0.1155 - acc:

0.9710 - val_loss: 0.1325 - val_acc: 0.9559
Epoch 22/30
186/187 [=====>.] - ETA: 0s - loss: 0.1127 - acc:
0.9714Epoch 1/30
188/187 [=====] - 6s 32ms/step - loss: 0.1123 - acc:
0.9717 - val_loss: 0.1333 - val_acc: 0.9591
Epoch 23/30
186/187 [=====>.] - ETA: 0s - loss: 0.1108 - acc:
0.9714Epoch 1/30
188/187 [=====] - 6s 32ms/step - loss: 0.1120 - acc:
0.9707 - val_loss: 0.1331 - val_acc: 0.9591
Epoch 24/30
186/187 [=====>.] - ETA: 0s - loss: 0.1117 - acc:
0.9717Epoch 1/30
188/187 [=====] - 6s 32ms/step - loss: 0.1117 - acc:
0.9714 - val_loss: 0.1324 - val_acc: 0.9559
Epoch 25/30
186/187 [=====>.] - ETA: 0s - loss: 0.1111 - acc:
0.9711Epoch 1/30
188/187 [=====] - 6s 32ms/step - loss: 0.1114 - acc:
0.9710 - val_loss: 0.1318 - val_acc: 0.9559
Epoch 26/30
186/187 [=====>.] - ETA: 0s - loss: 0.1114 - acc:
0.9707Epoch 1/30
188/187 [=====] - 6s 32ms/step - loss: 0.1113 - acc:
0.9710 - val_loss: 0.1314 - val_acc: 0.9559
Epoch 27/30
186/187 [=====>.] - ETA: 0s - loss: 0.1093 - acc:
0.9727Epoch 1/30
188/187 [=====] - 6s 32ms/step - loss: 0.1109 - acc:
0.9717 - val_loss: 0.1318 - val_acc: 0.9559
Epoch 28/30
186/187 [=====>.] - ETA: 0s - loss: 0.1108 - acc:
0.9721Epoch 1/30
188/187 [=====] - 6s 32ms/step - loss: 0.1106 - acc:
0.9720 - val_loss: 0.1307 - val_acc: 0.9575
Epoch 29/30
186/187 [=====>.] - ETA: 0s - loss: 0.1104 - acc:
0.9711Epoch 1/30
188/187 [=====] - 6s 32ms/step - loss: 0.1104 - acc:
0.9710 - val_loss: 0.1322 - val_acc: 0.9575
Epoch 30/30
186/187 [=====>.] - ETA: 0s - loss: 0.1096 - acc:
0.9738Epoch 1/30
188/187 [=====] - 6s 32ms/step - loss: 0.1101 - acc:
0.9734 - val_loss: 0.1305 - val_acc: 0.9559

Saving the model

```
model.save("VGG16_plant_deseas.h5")  
print("Saved model to disk")
```

Saved model to disk

Loading the model

```
model = models.load_model('VGG16_plant_deseas.h5')  
print("Model is loaded")
```

Model is loaded

Saving the Weights


```
model.save_weights('cnn_classification.h5')
```

Loading the weights


```
model.load_weights('cnn_classification.h5')
```

Visualization of training over epoch

```
train_acc = history.history['acc']  
val_acc = history.history['val_acc']  
train_loss = history.history['loss']  
val_loss = history.history['val_loss']  
  
epochs = range(len(train_acc))  
plt.plot(epochs, train_acc, 'b', label='Training Accuracy')  
plt.plot(epochs, val_acc, 'g', label='Validation Accuracy')  
plt.title('Training and Validation Accuracy')  
plt.grid()  
plt.legend()  
plt.figure()  
plt.show()  
  
plt.plot(epochs, train_loss, 'b', label='Training Loss')  
plt.plot(epochs, val_loss, 'g', label='Validation Loss')  
plt.title('Training and Validation Loss')  
plt.grid()  
plt.legend()  
plt.show()
```

<Figure size 640x480 with 0 Axes>



Performance measure

Get the filenames from the generator

```
fnames = test_generator.filenames
```

Get the ground truth from generator

```
ground_truth = test_generator.classes
```

Get the label to class mapping from the generator

```
label2index = test_generator.class_indices
```

Getting the mapping from class index to class label

```
idx2label = dict((v,k) for k,v in label2index.items())
```

Get the predictions from the model using the generator

```
predictions = model.predict_generator(test_generator,
steps=test_generator.samples/test_generator.batch_size,verbose=1)
predicted_classes = np.argmax(predictions,axis=1)
```

```
errors = np.where(predicted_classes != ground_truth)[0]
```

```
print("No of errors = {}/{}".format(len(errors),test_generator.samples))
```

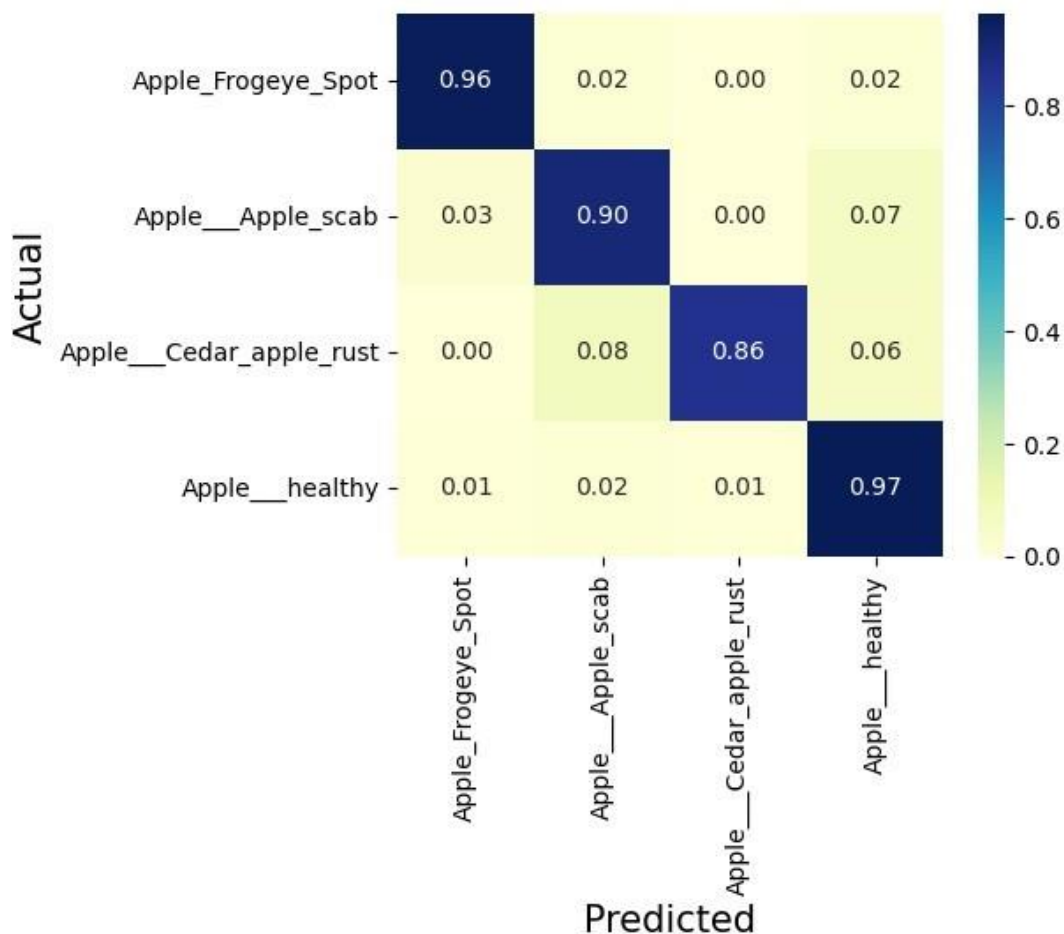
```
547/547 [=====] - 3s 6ms/step
```

```
No of errors = 33/547
```

```
accuracy = ((test_generator.samples-len(errors))/test_generator.samples) *
100
accuracy
```

```
93.96709323583181
```

```
from sklearn.metrics import confusion_matrix
import seaborn as sns
import numpy as np
from matplotlib import pyplot as plt
cm = confusion_matrix(y_true=ground_truth, y_pred=predicted_classes)
cm = np.array(cm)
# Normalise
cmn = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
fig, ax = plt.subplots(figsize=(5,4))
sns.heatmap(cmn, annot=True, fmt='.2f', xticklabels=label2index,
yticklabels=label2index, cmap="YlGnBu")
plt.ylabel('Actual', fontsize=15)
plt.xlabel('Predicted', fontsize=15)
plt.show(block=False)
```



```
from sklearn.metrics import classification_report
print(classification_report(ground_truth, predicted_classes,
target_names=label2index))
```

	precision	recall	f1-score	support
Apple_Frogeye_Spot	0.94	0.96	0.95	104
Apple___Apple_scab	0.92	0.90	0.91	134
Apple___Cedar_apple_rust	0.95	0.86	0.90	49
Apple___healthy	0.95	0.97	0.96	260
accuracy			0.94	547
macro avg	0.94	0.92	0.93	547
weighted avg	0.94	0.94	0.94	547

InceptionNet

```
from tensorflow.keras import applications
```

```
## Loading InceptionV3 model
```

```
base_model = applications.InceptionV3(weights="imagenet", include_top=False,
input_shape= (128, 128, 3))
base_model.trainable = False ## Not trainable weights
```

```
base_model.summary()
```

```
Downloading data from https://github.com/fchollet/deep-learning-
models/releases/download/v0.5/inception_v3_weights_tf_dim_ordering_tf_kernels
_notop.h5
87916544/87910968 [=====] - 15s 0us/step
Model: "inception_v3"
```

Layer (type)	Output Shape	Param #	Connected to
input_3 (InputLayer)	[(None, 128, 128, 3)]	0	
conv2d (Conv2D) input_3[0][0]	(None, 63, 63, 32)	864	
batch_normalization (BatchNorma	(None, 63, 63, 32)	96	conv2d[0][0]

activation (Activation) batch_normalization[0][0]	(None, 63, 63, 32)	0
conv2d_1 (Conv2D) activation[0][0]	(None, 61, 61, 32)	9216
batch_normalization_1 (BatchNor conv2d_1[0][0]	(None, 61, 61, 32)	96
activation_1 (Activation) batch_normalization_1[0][0]	(None, 61, 61, 32)	0
conv2d_2 (Conv2D) activation_1[0][0]	(None, 61, 61, 64)	18432
batch_normalization_2 (BatchNor conv2d_2[0][0]	(None, 61, 61, 64)	192
activation_2 (Activation) batch_normalization_2[0][0]	(None, 61, 61, 64)	0
max_pooling2d (MaxPooling2D) activation_2[0][0]	(None, 30, 30, 64)	0
conv2d_3 (Conv2D) max_pooling2d[0][0]	(None, 30, 30, 80)	5120
batch_normalization_3 (BatchNor conv2d_3[0][0]	(None, 30, 30, 80)	240
activation_3 (Activation) batch_normalization_3[0][0]	(None, 30, 30, 80)	0
conv2d_4 (Conv2D) activation_3[0][0]	(None, 28, 28, 192)	138240
batch_normalization_4 (BatchNor	(None, 28, 28, 192)	576

conv2d_91[0][0]

batch_normalization_92 (Batch Normalization)	(None, 2, 2, 384)	1152
--	-------------------	------

conv2d_92[0][0]

conv2d_93 (Conv2D)	(None, 2, 2, 192)	393216
--------------------	-------------------	--------

average_pooling2d_8[0][0]

batch_normalization_85 (Batch Normalization)	(None, 2, 2, 320)	960
--	-------------------	-----

conv2d_85[0][0]

activation_87 (Activation)	(None, 2, 2, 384)	0
----------------------------	-------------------	---

batch_normalization_87[0][0]

activation_88 (Activation)	(None, 2, 2, 384)	0
----------------------------	-------------------	---

batch_normalization_88[0][0]

activation_91 (Activation)	(None, 2, 2, 384)	0
----------------------------	-------------------	---

batch_normalization_91[0][0]

activation_92 (Activation)	(None, 2, 2, 384)	0
----------------------------	-------------------	---

batch_normalization_92[0][0]

batch_normalization_93 (Batch Normalization)	(None, 2, 2, 192)	576
--	-------------------	-----

conv2d_93[0][0]

activation_85 (Activation)	(None, 2, 2, 320)	0
----------------------------	-------------------	---

batch_normalization_85[0][0]

mixed9_1 (Concatenate)	(None, 2, 2, 768)	0
------------------------	-------------------	---

activation_87[0][0]

activation_88[0][0]

concatenate_1 (Concatenate)	(None, 2, 2, 768)	0
-----------------------------	-------------------	---

activation_91[0][0]

activation_92[0][0]

activation_93 (Activation)	(None, 2, 2, 192)	0
batch_normalization_93[0][0]		

mixed10 (Concatenate)	(None, 2, 2, 2048)	0
activation_85[0][0]		

mixed9_1[0][0]

concatenate_1[0][0]

activation_93[0][0]

=====

Total params: 21,802,784

Trainable params: 0

Non-trainable params: 21,802,784

```
flatten_layer = layers.GlobalAveragePooling2D()
dense_layer_1 = layers.Dense(64, activation='relu')
dense_layer_2 = layers.Dense(32, activation='relu')
prediction_layer = layers.Dense(4, activation='softmax')
```

```
model = models.Sequential([
    base_model,
    flatten_layer,
    dense_layer_1,
    dense_layer_2,
    prediction_layer
])
```

```
model.summary()
```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
inception_v3 (Model)	(None, 2, 2, 2048)	21802784
global_average_pooling2d_2	(None, 2048)	0
dense_2 (Dense)	(None, 64)	131136
dense_3 (Dense)	(None, 32)	2080

dense_4 (Dense)	(None, 4)	132
-----------------	-----------	-----

=====
 Total params: 21,936,132
 Trainable params: 133,348
 Non-trainable params: 21,802,784

```

model.compile(optimizer = Adam(learning_rate = 0.001),
loss='categorical_crossentropy', metrics=['acc'])
# Train the model
history = model.fit(train_generator,
                    steps_per_epoch=train_generator.samples/train_generator.batch_size,
                    epochs=10,
                    validation_data=validation_generator,

validation_steps=validation_generator.samples/validation_generator.batch_size
,
                    verbose=2)

```

```

Epoch 1/10
Epoch 1/10
188/187 - 9s - loss: 0.3544 - acc: 0.8725 - val_loss: 1.8910 - val_acc:
0.7134
Epoch 2/10
Epoch 1/10
188/187 - 5s - loss: 0.4988 - acc: 0.8113 - val_loss: 1.5254 - val_acc:
0.7181
Epoch 3/10
Epoch 1/10
188/187 - 5s - loss: 0.4549 - acc: 0.8246 - val_loss: 1.2376 - val_acc:
0.7433
Epoch 4/10
Epoch 1/10
188/187 - 5s - loss: 0.4242 - acc: 0.8342 - val_loss: 1.2466 - val_acc:
0.7780
Epoch 5/10
Epoch 1/10
188/187 - 5s - loss: 0.4462 - acc: 0.8382 - val_loss: 1.6309 - val_acc:
0.7165
Epoch 6/10
Epoch 1/10
188/187 - 5s - loss: 0.4180 - acc: 0.8445 - val_loss: 1.3079 - val_acc:
0.7260
Epoch 7/10
Epoch 1/10
188/187 - 5s - loss: 0.4094 - acc: 0.8442 - val_loss: 2.0786 - val_acc:
0.6693
Epoch 8/10
Epoch 1/10

```



```
188/187 - 5s - loss: 0.3816 - acc: 0.8595 - val_loss: 1.5253 - val_acc: 0.7606
```

```
Epoch 9/10
```

```
Epoch 1/10
```

```
188/187 - 5s - loss: 0.4081 - acc: 0.8449 - val_loss: 1.1693 - val_acc: 0.7606
```

```
Epoch 10/10
```

```
Epoch 1/10
```

```
188/187 - 5s - loss: 0.3675 - acc: 0.8615 - val_loss: 1.4619 - val_acc: 0.7276
```

```
model.save("InceptionNet_plant_deseas.h5")
```

```
print("Saved model to disk")
```

```
Saved model to disk
```

```
train_acc = history.history['acc']
```

```
val_acc = history.history['val_acc']
```

```
train_loss = history.history['loss']
```

```
val_loss = history.history['val_loss']
```

```
epochs = range(len(train_acc))
```

```
plt.plot(epochs, train_acc, 'b', label='Training Accuracy')
```

```
plt.plot(epochs, val_acc, 'g', label='Validation Accuracy')
```

```
plt.title('Training and Validation Accuracy')
```

```
plt.grid()
```

```
plt.legend()
```

```
plt.figure()
```

```
plt.show()
```

```
plt.plot(epochs, train_loss, 'b', label='Training Loss')
```


```
plt.plot(epochs, val_loss, 'g', label='Validation Loss')
```

```
plt.title('Training and Validation Loss')
```


```
plt.grid()
```

```
plt.legend()
```

```
plt.show()
```



<Figure size 640x480 with 0 Axes>



```
# Get the filenames from the generator
fnames = test_generator.filenames

# Get the ground truth from generator
ground_truth = test_generator.classes

# Get the Label to class mapping from the generator
label2index = test_generator.class_indices

# Getting the mapping from class index to class label
idx2label = dict((v,k) for k,v in label2index.items())

# Get the predictions from the model using the generator
predictions = model.predict_generator(test_generator,
steps=test_generator.samples/test_generator.batch_size,verbose=1)
predicted_classes = np.argmax(predictions,axis=1)

errors = np.where(predicted_classes != ground_truth)[0]
print("No of errors = {}/{}".format(len(errors),test_generator.samples))

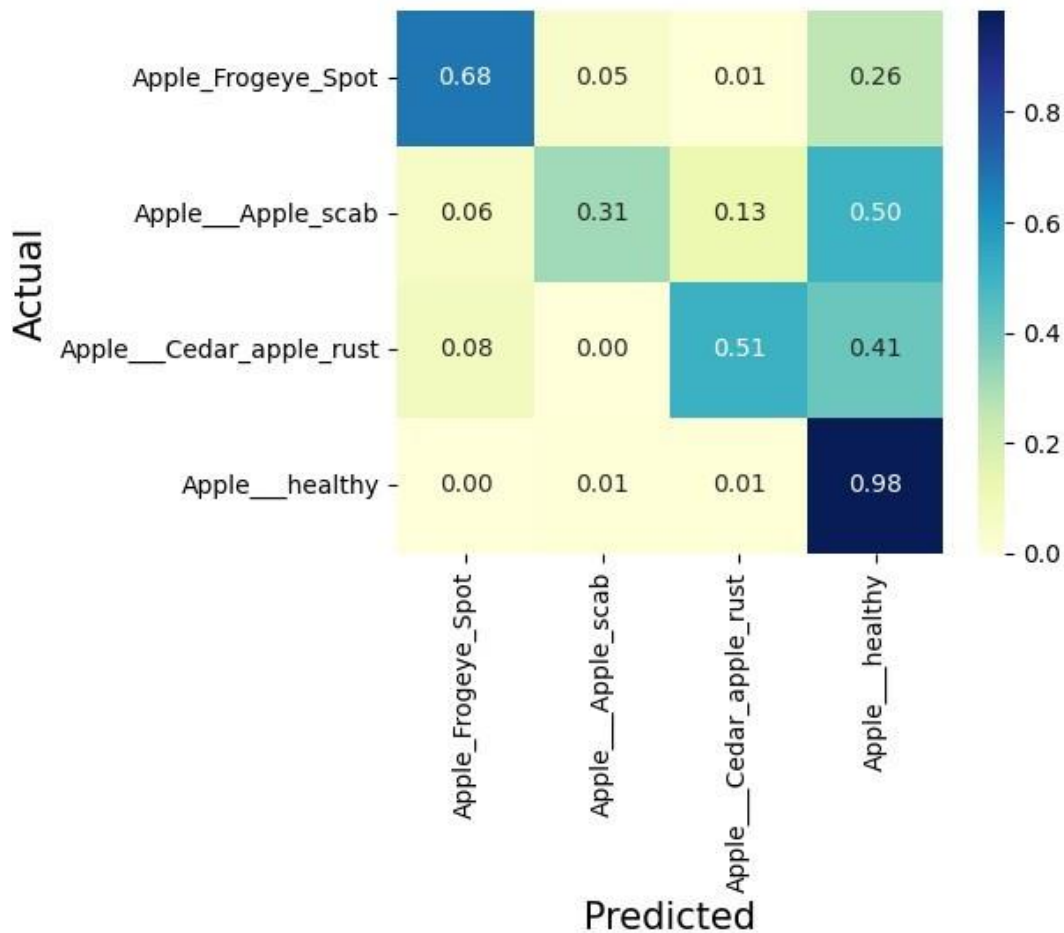
547/547 [=====] - 9s 17ms/step
No of errors = 153/547
```

```
accuracy = ((test_generator.samples-len(errors))/test_generator.samples) *
100
```

```
accuracy
```

```
72.0292504570384
```

```
from sklearn.metrics import confusion_matrix
import seaborn as sns
import numpy as np
from matplotlib import pyplot as plt
cm = confusion_matrix(y_true=ground_truth, y_pred=predicted_classes)
cm = np.array(cm)
# Normalise
cmn = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
fig, ax = plt.subplots(figsize=(5,4))
sns.heatmap(cmn, annot=True, fmt='.2f', xticklabels=label2index,
yticklabels=label2index, cmap="YlGnBu")
plt.ylabel('Actual', fontsize=15)
plt.xlabel('Predicted', fontsize=15)
plt.show(block=False)
```



```
from sklearn.metrics import classification_report
print(classification_report(ground_truth, predicted_classes,
target_names=label2index))
```

	precision	recall	f1-score	support
Apple_Frogeye_Spot	0.86	0.68	0.76	104
Apple___Apple_scab	0.86	0.31	0.46	134
Apple___Cedar_apple_rust	0.56	0.51	0.53	49
Apple___healthy	0.69	0.98	0.81	260
accuracy			0.72	547
macro avg	0.74	0.62	0.64	547
weighted avg	0.75	0.72	0.69	547

ResNet

```
from keras import applications
```

```
## Loading VGG16 model
```

```
base_model = applications.ResNet50(weights="imagenet", include_top=False,
input_shape= (128, 128, 3))
base_model.trainable = False ## Not trainable weights
```

```
base_model.summary()
```

```
flatten_layer = layers.GlobalAveragePooling2D()
# dense_layer_1 = layers.Dense(63, activation='relu')
# dense_layer_2 = layers.Dense(32, activation='relu')
prediction_layer = layers.Dense(4, activation='softmax')
```

```
model = models.Sequential([
    base_model,
    flatten_layer,
    # dense_layer_1,
    # dense_layer_2,
    prediction_layer
])
```

```
model.summary()
```

```
model.compile(optimizer = Adam(learning_rate = 0.001),
loss='categorical_crossentropy', metrics=['acc'])
# Train the model
history = model.fit(train_generator,
    steps_per_epoch=train_generator.samples/train_generator.batch_size,
    epochs=30,
```

```

        validation_data=validation_generator,

validation_steps=validation_generator.samples/validation_generator.batch_size
',
        verbose=1)

model.save("ResNet_plant_deseas.h5")
print("Saved model to disk")

train_acc = history.history['acc']
val_acc = history.history['val_acc']
train_loss = history.history['loss']
val_loss = history.history['val_loss']

epochs = range(len(train_acc))
plt.plot(epochs, train_acc, 'b', label='Training Accuracy')
plt.plot(epochs, val_acc, 'g', label='Validation Accuracy')
plt.title('Training and Validation Accuracy')
plt.grid()
plt.legend()
plt.figure()
plt.show()

plt.plot(epochs, train_loss, 'b', label='Training Loss')
plt.plot(epochs, val_loss, 'g', label='Validation Loss')
plt.title('Training and Validation Loss')
plt.grid()
plt.legend()
plt.figure()
plt.show()

# Get the filenames from the generator
fnames = test_generator.filenames

# Get the ground truth from generator
ground_truth = test_generator.classes

# Get the label to class mapping from the generator
label2index = test_generator.class_indices

# Getting the mapping from class index to class label
idx2label = dict((v,k) for k,v in label2index.items())

# Get the predictions from the model using the generator
predictions = model.predict_generator(test_generator,
steps=test_generator.samples/test_generator.batch_size,verbose=1)
predicted_classes = np.argmax(predictions,axis=1)

errors = np.where(predicted_classes != ground_truth)[0]
print("No of errors = {}/{}".format(len(errors),test_generator.samples))

```

```

accuracy = ((test_generator.samples-len(errors))/test_generator.samples) *
100
accuracy

from sklearn.metrics import confusion_matrix
import seaborn as sns
import numpy as np
from matplotlib import pyplot as plt
cm = confusion_matrix(y_true=ground_truth, y_pred=predicted_classes)
cm = np.array(cm)
# Normalise
cmn = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
fig, ax = plt.subplots(figsize=(5,4))
sns.heatmap(cmn, annot=True, fmt='.2f', xticklabels=label2index,
yticklabels=label2index, cmap="YlGnBu")
plt.ylabel('Actual', fontsize=15)
plt.xlabel('Predicted', fontsize=15)
plt.show(block=False)

from sklearn.metrics import classification_report
print(classification_report(ground_truth, predicted_classes,
target_names=label2index))

```

Experiment - 5

Text data handling with RNN for sentiment analysis

Importing Necessary libraries

```
import pandas as pd      # to Load dataset
import numpy as np       # for mathematic equation
from nltk.corpus import stopwords # to get collection of stopwords
from sklearn.model_selection import train_test_split # for splitting dataset
from tensorflow.keras.preprocessing.text import Tokenizer # to encode text to int
from tensorflow.keras.preprocessing.sequence import pad_sequences # to do padding or truncating
from tensorflow.keras.models import Sequential # the model
from tensorflow.keras.layers import Embedding, LSTM, Dense # layers of the architecture
from tensorflow.keras.callbacks import ModelCheckpoint # save model
from tensorflow.keras.models import load_model # Load saved model
import re
from keras.layers import SimpleRNN
```

Preparing the data named IMDB

```
data = pd.read_csv('/content/ IMDB Dataset.csv')
print(data)
```

	review	sentiment
0	One of the other reviewers has mentioned that ...	positive
1	A wonderful little production. The...	positive
2	I thought this was a wonderful way to spend ti...	positive
3	Basically there's a family where a little boy ...	negative
4	Petter Mattei's "Love in the Time of Money" is...	positive
...
49995	I thought this movie did a down right good job...	positive
49996	Bad plot, bad dialogue, bad acting, idiotic di...	negative
49997	I am a Catholic taught in parochial elementary...	negative
49998	I'm going to have to disagree with the previou...	negative
49999	No one expects the Star Trek movies to be high...	negative

```
[50000 rows x 2 columns]
```



```

import nltk
nltk.download("stopwords")
english_stops = set(stopwords.words('english'))

[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data]   Unzipping corpora/stopwords.zip.

def load_dataset():
    df = pd.read_csv('/content/ IMDB Dataset.csv')
    x_data = df['review']      # Reviews/Input
    y_data = df['sentiment']   # Sentiment/Output

    # PRE-PROCESS REVIEW
    x_data = x_data.replace({'<.*?>': ''}, regex = True)      # remove
    # html tag
    x_data = x_data.replace({'[^A-Za-z]': ' '}, regex = True) # remove
    # non alphabet
    x_data = x_data.apply(lambda review: [w for w in review.split() if w not
in english_stops]) # remove stop words
    x_data = x_data.apply(lambda review: [w.lower() for w in review]) #
    # lower case

    # ENCODE SENTIMENT -> 0 & 1
    y_data = y_data.replace('positive', 1)
    y_data = y_data.replace('negative', 0)

    return x_data, y_data

x_data, y_data = load_dataset()

print('Reviews')
print(x_data, '\n')
print('Sentiment')
print(y_data)

```

Reviews

```

0      [one, reviewers, mentioned, watching, oz, epis...
1      [a, wonderful, little, production, the, filmin...
2      [i, thought, wonderful, way, spend, time, hot,...
3      [basically, family, little, boy, jake, thinks,...
4      [petter, mattei, love, time, money, visually, ...
...
49995  [i, thought, movie, right, good, job, it, crea...
49996  [bad, plot, bad, dialogue, bad, acting, idioti...
49997  [i, catholic, taught, parochial, elementary, s...
49998  [i, going, disagree, previous, comment, side, ...
49999  [no, one, expects, star, trek, movies, high, a...
Name: review, Length: 50000, dtype: object

```

```

Sentiment
0      1
1      1
2      1
3      0
4      1
..
49995   1
49996   0
49997   0
49998   0
49999   0
Name: sentiment, Length: 50000, dtype: int64

```

Split Dataset

```

x_train, x_test, y_train, y_test = train_test_split(x_data, y_data, test_size
= 0.2)

```

```

print('Train Set')
print(x_train, '\n')
print(x_test, '\n')
print('Test Set')
print(y_train, '\n')
print(y_test)

```

```

Train Set
22596 [boring, badly, written, italian, exploitation...
5353  [the, other, supposed, horror, movie, made, it...
42152 [a, tough, life, gets, tougher, three, childre...
15434 [why, earth, colin, firth, pointless, film, ha...
7280  [this, far, worst, movie, i, ever, seen, cinem...
...
39945 [this, show, lasted, moments, plots, usually, ...
13858 [i, rented, thinking, would, pretty, good, cov...
25266 [having, pleasantly, surprised, sandra, bulloc...
10659 [the, difficulty, i, musical, version, les, mi...
39372 [this, movie, proof, film, noire, enduring, st...
Name: review, Length: 40000, dtype: object

```

```

2006  [this, movie, time, favorite, you, really, see...
33575 [this, british, film, version, stage, play, i,...
6808  [alexander, nevsky, brilliant, piece, cinemati...
32330 [found, old, vhs, version, film, parents, hous...
3777  [i, went, see, movie, daughter, i, insisted, g...
...
40255 [what, heck, people, expect, horror, films, da...

```

```
5864    [especially, time, much, science, fiction, fil...
44604   [nicole, eggert, listed, star, despite, michea...
42481   [a, thief, night, got, best, end, times, thril...
31671   [i, enjoy, national, anthem, i, enjoy, nationa...
Name: review, Length: 10000, dtype: object
```

```
Test Set
22596    0
5353     0
42152    1
15434    0
7280     0
..
39945    0
13858    0
25266    0
10659    0
39372    1
Name: sentiment, Length: 40000, dtype: int64
```

```
2006     1
33575    1
6808     1
32330    0
3777     0
..
40255    1
5864     1
44604    0
42481    1
31671    1
Name: sentiment, Length: 10000, dtype: int64
```

```
def get_max_length():
    review_length = []
    for review in x_train:
        review_length.append(len(review))

    return int(np.ceil(np.mean(review_length)))

# ENCODE REVIEW
token = Tokenizer(lower=False)    # no need lower, because already Lowered
the data in Load_data()
token.fit_on_texts(x_train)
x_train = token.texts_to_sequences(x_train)
x_test = token.texts_to_sequences(x_test)

max_length = get_max_length()
```

```

x_train = pad_sequences(x_train, maxlen=max_length, padding='post',
truncating='post')
x_test = pad_sequences(x_test, maxlen=max_length, padding='post',
truncating='post')

total_words = len(token.word_index) + 1    # add 1 because of 0 padding
print('Total Words:', total_words)

print('Encoded X Train\n', x_train, '\n')
print('Encoded X Test\n', x_test, '\n')
print('Maximum review length: ', max_length)

```

Total Words: 92636

Encoded X Train

```

[[ 257  863  310 ...    0    0    0]
 [   2 1340  350 ...   28  282 409]
 [  39 1138   40 ...    0    0    0]
 ...
 [1587 3903  660 ...   62 14457 1006]
 [   2 6090    1 ...  4973  5675  406]
 [   8    3 2912 ...    0    0    0]]

```

Encoded X Test

```

[[   8    3  10 ...    0    0    0]
 [   8  603    4 ...  278 10278 2289]
 [3551 11276  417 ...    0    0    0]
 ...
 [3908 20405 3718 ...    0    0    0]
 [  39 2984  218 ...  3947    3  765]
 [   1  260 1833 ...    0    0    0]]

```

Maximum review length: 130

Build Architecture/Model Embedding Layer

```

rnn = Sequential()

rnn.add(Embedding(total_words,32,input_length =max_length))
rnn.add(SimpleRNN(64,input_shape = (total_words, max_length),
return_sequences=False,activation="relu"))
rnn.add(Dense(1, activation = 'sigmoid')) #flatten

print(rnn.summary())
rnn.compile(loss="binary_crossentropy",optimizer='adam',metrics=["accuracy"])

```

Model: "sequential"

Layer (type)	Output Shape	Param #
embedding (Embedding)	(None, 130, 32)	2964352
simple_rnn (SimpleRNN)	(None, 64)	6208
dense (Dense)	(None, 1)	65

=====
Total params: 2,970,625
Trainable params: 2,970,625
Non-trainable params: 0

None

#Trainin the Model

```
history = rnn.fit(x_train,y_train,epochs = 20,batch_size=128,verbose = 1)
```

Epoch 1/20

313/313 [=====] - 96s 286ms/step - loss: 0.6915 - accuracy: 0.5184

Epoch 2/20

313/313 [=====] - 68s 218ms/step - loss: 0.6616 - accuracy: 0.5879

Epoch 3/20

313/313 [=====] - 63s 202ms/step - loss: 0.6626 - accuracy: 0.5762

Epoch 4/20

313/313 [=====] - 56s 180ms/step - loss: 0.5900 - accuracy: 0.6328

Epoch 5/20

313/313 [=====] - 51s 162ms/step - loss: 0.4166 - accuracy: 0.8135

Epoch 6/20

313/313 [=====] - 51s 161ms/step - loss: 0.3511 - accuracy: 0.8806

Epoch 7/20

313/313 [=====] - 47s 149ms/step - loss: 0.2362 - accuracy: 0.9194

Epoch 8/20

313/313 [=====] - 46s 148ms/step - loss: 0.1676 - accuracy: 0.9421

Epoch 9/20

313/313 [=====] - 45s 144ms/step - loss: 0.3168 - accuracy: 0.8689

Epoch 10/20

313/313 [=====] - 46s 148ms/step - loss: 0.5571 -

```

accuracy: 0.6458
Epoch 11/20
313/313 [=====] - 44s 141ms/step - loss: 0.4791 -
accuracy: 0.7574
Epoch 12/20
313/313 [=====] - 45s 144ms/step - loss: 0.2817 -
accuracy: 0.9088
Epoch 13/20
313/313 [=====] - 43s 139ms/step - loss: 0.4030 -
accuracy: 0.8431
Epoch 14/20
313/313 [=====] - 45s 142ms/step - loss: 0.2630 -
accuracy: 0.9154
Epoch 15/20
313/313 [=====] - 43s 138ms/step - loss: 0.2351 -
accuracy: 0.9260
Epoch 16/20
313/313 [=====] - 43s 138ms/step - loss: 0.2000 -
accuracy: 0.9385
Epoch 17/20
313/313 [=====] - 45s 143ms/step - loss: 0.1599 -
accuracy: 0.9498
Epoch 18/20
313/313 [=====] - 44s 140ms/step - loss: 0.1374 -
accuracy: 0.9577
Epoch 19/20
313/313 [=====] - 44s 141ms/step - loss: 0.1331 -
accuracy: 0.9611
Epoch 20/20
313/313 [=====] - 43s 136ms/step - loss: 0.2814 -
accuracy: 0.8869

```

Saving The Model

```

model = rnn.save('rnn.h5')
loaded_model = load_model('rnn.h5')

```

Evaluation

```

y_pred = rnn.predict(x_test, batch_size = 128)
print(y_pred)
print(y_test)
for i in range(len(y_pred)):
    if y_pred[i]>0.5:
        y_pred[i] = 1
    else:
        y_pred[i] = 0

true = 0
for i, y in enumerate(y_test):
    if y == y_pred[i]:

```

```

true += 1

print('Correct Prediction: {}'.format(true))
print('Wrong Prediction: {}'.format(len(y_pred) - true))
print('Accuracy: {}'.format(true/len(y_pred)*100))

79/79 [=====] - 1s 12ms/step
[[0.78446704]
 [0.02569966]
 [0.78301245]
 ...
 [0.2700789 ]
 [0.72713566]
 [0.78446704]]
2006      1
33575     1
6808      1
32330     0
3777      0
..
40255     1
5864      1
44604     0
42481     1
31671     1
Name: sentiment, Length: 10000, dtype: int64
Correct Prediction: 6918
Wrong Prediction: 3082
Accuracy: 69.17999999999999

```

Message: Nothing was typical about this. Everything was beautifully done in this movie, the story, the flow, the scenario, everything. I highly recommend it for mystery lovers, for anyone who wants to watch a good movie!

Example review

```
review = str(input('Movie Review: '))
```

Movie Review: Nothing was typical about this. Everything was beautifully done in this movie, the story, the flow, the scenario, everything. I highly recommend it for mystery lovers, for anyone who wants to watch a good movie!

#Pre-processing of entered review

```

# Pre-process input
regex = re.compile(r'^a-zA-Z\s')
review = regex.sub('', review)
print('Cleaned: ', review)

```

```

words = review.split(' ')
filtered = [w for w in words if w not in english_stops]
filtered = ' '.join(filtered)
filtered = [filtered.lower()]

```

```

print('Filtered: ', filtered)

```

Cleaned: Nothing was typical about this Everything was beautifully done in this movie the story the flow the scenario everything I highly recommend it for mystery lovers for anyone who wants to watch a good movie

Filtered: ['nothing typical everything beautifully done movie story flow scenario everything i highly recommend mystery lovers anyone wants watch good movie']

```

tokenize_words = token.texts_to_sequences(filtered)
tokenize_words = pad_sequences(tokenize_words, maxlen=max_length,
padding='post', truncating='post')
print(tokenize_words)

```

```

[[ 76 705 174 1210 126 3 13 2692 2596 174 1 442 280 701
 1771 155 400 33 9 3 0 0 0 0 0 0 0 0
 0 0 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 0 0]]

```

#Prediction

```

result = rnn.predict(tokenize_words)
print(result)

```

```

1/1 [=====] - 0s 40ms/step
[[0.78446704]]

```

```

if result >= 0.7:
    print('positive')
else:
    print('negative')

```

positive

Experiment - 6

Sentiment analysis using RNN-LSTM on tweets data

Importing Necessary libraries

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

Preparing the data named IMDB

```
df = pd.read_csv('/content/data.csv')
df.head()
```

	Unnamed: 0	count	hate_speech	offensive_language	neither	class	\
0	0	3	0	0	3	2	
1	1	3	0	3	0	1	
2	2	3	0	3	0	1	
3	3	3	0	2	1	1	
4	4	6	0	6	0	1	

```
                                tweet
0  !!! RT @mayasolovely: As a woman you shouldn't...
1  !!!!!! RT @mleew17: boy dats cold...tyga dwn ba...
2  !!!!!! RT @UrKindOfBrand Dawg!!!! RT @80sbaby...
3  !!!!!! RT @C_G_Anderson: @viva_based she lo...
4  !!!!!! RT @ShenikaRoberts: The shit you...
```

```
classes = ['Hate Speech', 'Offensive Language', 'None']
```

```
df.drop(['count', 'hate_speech', 'offensive_language', 'neither', 'Unnamed: 0'], axis=1, inplace=True)
```

```
df.head()
```

	class	tweet
0	2	!!! RT @mayasolovely: As a woman you shouldn't...
1	1	!!!!!! RT @mleew17: boy dats cold...tyga dwn ba...
2	1	!!!!!! RT @UrKindOfBrand Dawg!!!! RT @80sbaby...
3	1	!!!!!! RT @C_G_Anderson: @viva_based she lo...
4	1	!!!!!! RT @ShenikaRoberts: The shit you...


```
df.shape
```

```
(24783, 2)
```

Statistics of the Data

```
labels = df['class']
unique, counts = np.unique(labels, return_counts=True)
values = list(zip(unique, counts))
plt.bar(classes, counts)
for i in values:
    print(classes[i[0]], ' : ', i[1])
plt.show()
```

```
Hate Speech : 1430
Offensive Language : 19190
None : 4163
```



```
hate_tweets = df[df['class']==0]
offensive_tweets = df[df['class']==1]
neither = df[df['class']==2]
print(hate_tweets.shape)
print(offensive_tweets.shape)
print(neither.shape)

(1430, 2)
(19190, 2)
(4163, 2)

for i in range(3):
    hate_tweets = pd.concat([hate_tweets, hate_tweets], ignore_index = True)
```

```

neither = pd.concat([neither,neither,neither], ignore_index = True)
offensive_tweets = offensive_tweets.iloc[0:12000,:]
print(hate_tweets.shape)
print(offensive_tweets.shape)
print(neither.shape)

(11440, 2)
(12000, 2)
(12489, 2)


df = pd.concat([hate_tweets,offensive_tweets,neither],ignore_index = True)
df.shape

(35929, 2)

labels = df['class']
unique, counts = np.unique(labels, return_counts=True)
values = list(zip(unique, counts))
plt.bar(classes,counts)
for i in values:
    print(classes[i[0]], ' : ',i[1])
plt.show()

Hate Speech : 11440
Offensive Language : 12000
None : 12489

```



```
df.head()
```

```
class tweet
0      0 "@Blackman38Tide: @WhaleLookyHere @HowdyDowdy1...
1      0 "@CB_Baby24: @white_thunduh alsarabsss" hes a ...
2      0 "@DevilGrimz: @VigxRArts you're fucking gay, b...
3      0 "@MarkRoundtreeJr: LMFA0000 I HATE BLACK PEOP...
4      0 "@NoChillPaz: "At least I'm not a nigger" http...
```

Preprocessing

```
import nltk
from nltk.stem import WordNetLemmatizer
from nltk.corpus import stopwords
import re
nltk.download('wordnet')
nltk.download('stopwords')

[nltk_data] Downloading package wordnet to /root/nltk_data...
[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data] Unzipping corpora/stopwords.zip.
```

```
True
```

```
# dealing with Slangs
```

```
d = {'luv': 'love', 'wud': 'would', 'lyk': 'like', 'wateva': 'whatever', 'ttyl': 'talk
to you later',
      'kul': 'cool', 'fyn': 'fine', 'omg': 'oh my
god!', 'fam': 'family', 'bruh': 'brother',
      'cud': 'could', 'fud': 'food', 'u': 'you',
      'ur': 'your', 'bday' : 'birthday', 'bihday' : 'birthday'}
```

```
stop_words = set(stopwords.words("english"))
stop_words.add('rt')
stop_words.remove('not')
lemmatizer = WordNetLemmatizer()
giant_url_regex = ('http[s]?://(?:[a-zA-Z]|[0-9]|[$-_@.&+]|'
'[*\(\),]|(?:%[0-9a-fA-F][0-9a-fA-F]))+')
mention_regex = '@[\w\-\]+'
```

```
def clean_text(text):
    text = re.sub('\"', "", text)
    text = re.sub(mention_regex, ' ', text) #removing all user names
    text = re.sub(giant_url_regex, ' ', text) #removing the urls
    text = text.lower()
    text = re.sub("hm+", "", text) #removing variants of hmmm
    text = re.sub("[^a-z]+", " ", text) #removing all numbers, special chars
like @, #, ? etc
    text = text.split()
    text = [word for word in text if not word in stop_words]
    text = [d[word] if word in d else word for word in text] #replacing some
```

slangs

```
text = [lemmatizer.lemmatize(token) for token in text]
text = [lemmatizer.lemmatize(token, "v") for token in text]
text = " ".join(text)
return text
```

```
df['processed_tweets'] = df.tweet.apply(lambda x: clean_text(x)) #
df.review.map(clean_text) Also can be used
df.head()
```

```
class tweet \
0      0 "@Blackman38Tide: @WhaleLookyHere @HowdyDowdy1...
1      0 "@CB_Baby24: @white_thunduh alсарabsss" hes a ...
2      0 "@DevilGrimz: @VigxRArts you're fucking gay, b...
3      0 "@MarkRoundtreeJr: LMFAOOOO I HATE BLACK PEOP...
4      0 "@NoChillPaz: "At least I'm not a nigger" http...
```

```
processed_tweets
0      queer gaywad
1      alсарabsss he beaner smh tell he mexican
2      fuck gay blacklist hoe hold tehgodclan anyway
3      lmfaoooo hate black people black people nigger
4      least not nigger lmfao
```

```
x = df.processed_tweets
y = df['class']
print(x.shape)
print(y.shape)

(35929,)
(35929,)
```

finding unique words

```
word_unique = []
for i in x:
    for j in i.split():
        word_unique.append(j)
unique, counts = np.unique(word_unique, return_counts=True)
print("The total words in the tweets are : ", len(word_unique))
print("The total UNIQUE words in the tweets are : ", len(unique))
```

The total words in the tweets are : 275540

The total UNIQUE words in the tweets are : 14146

finding length of tweets

```
tweets_length = []
for i in x:
    tweets_length.append(len(i.split()))
print("The Average Length tweets are : ", np.mean(tweets_length))
print("The max length of tweets is : ", np.max(tweets_length))
print("The min length of tweets is : ", np.min(tweets_length))
```

The Average Length tweets are : 7.669013888502324
The max length of tweets is : 28
The min length of tweets is : 0

```
tweets_length = pd.DataFrame(tweets_length)
# tweets_length.describe()
```

```

count    35929.000000
mean      7.669014
std       3.989625
min       0.000000
25%       4.000000
50%       7.000000
75%      11.000000
max      28.000000
```

```
#Sorting the Unique words based on their Frequency
col = list(zip(unique, counts))
col = sorted(col, key = lambda x: x[1],reverse=True)
col=pd.DataFrame(col)
print("Top 20 Occuring Words with their frequency are:")
col.iloc[:20,:]
```

Top 20 Occuring Words with their frequency are:

	0	1
0	bitch	9066
1	like	3817
2	get	3636
3	hoe	3426
4	trash	3217
5	fuck	3103
6	nigga	2819
7	faggot	2239
8	as	2073
9	you	1851
10	pussy	1847
11	go	1773
12	not	1764
13	bird	1515
14	lol	1494
15	nigger	1459
16	say	1456
17	make	1373
18	amp	1329
19	white	1328

```
from sklearn.feature_extraction.text import TfidfVectorizer
```

```

vectorizer = TfidfVectorizer(max_features = 8000 )
# tokenize and build vocab

vectorizer.fit(x)
# summarize

print(len(vectorizer.vocabulary_))
print(vectorizer.idf_.shape)

8000
(8000,)

x_tfidf = vectorizer.transform(x).toarray()
print(x_tfidf.shape)

(35929, 8000)

from keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences

num_words = 8000
embed_dim = 32
tokenizer = Tokenizer(num_words=num_words,oov_token = "<oov>" )
tokenizer.fit_on_texts(x)
word_index=tokenizer.word_index
sequences = tokenizer.texts_to_sequences(x)
length=[]
for i in sequences:
    length.append(len(i))
print(len(length))
print("Mean is: ",np.mean(length))
print("Max is: ",np.max(length))
print("Min is: ",np.min(length))

35929
Mean is:  7.669013888502324
Max is:  28
Min is:  0

pad_length = 24
sequences = pad_sequences(sequences, maxlen = pad_length, truncating = 'pre',
padding = 'post')
sequences.shape

(35929, 24)

#Splitting the Data

from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test = train_test_split(sequences,y,test_size =
0.05)
print(x_train.shape)

```

```
print(x_test.shape)
print(y_train.shape)
print(y_test.shape)
```

```
(34132, 24)
(1797, 24)
(34132,)
(1797,)
```

#RNN Model

```
from keras.layers import Dense, Embedding, Dropout, Activation, Flatten,
SimpleRNN
from keras.layers import GlobalMaxPool1D
from keras.models import Model, Sequential
import tensorflow as tf
```

```
recall = tf.keras.metrics.Recall()
precision = tf.keras.metrics.Precision()
```

```
model = Sequential([Embedding(num_words, embed_dim, input_length =
pad_length),
```

```
                    SimpleRNN(8, return_sequences = True),
                    GlobalMaxPool1D(),
                    Dense(20, activation =
'relu', kernel_initializer='he_uniform'),
                    Dropout(0.25),
                    Dense(3, activation = 'softmax')])
```

```
model.compile(loss='sparse_categorical_crossentropy', optimizer='adam',
metrics=['accuracy'])
```

```
# model.name = 'Twitter Hate Text Classification'
```

```
model.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
=====		
embedding_1 (Embedding)	(None, 24, 32)	256000
simple_rnn_1 (SimpleRNN)	(None, 24, 8)	328
global_max_pooling1d_1 (GlobalMaxPooling1D)	(None, 8)	0
dense_2 (Dense)	(None, 20)	180
dropout_1 (Dropout)	(None, 20)	0
dense_3 (Dense)	(None, 3)	63


```
=====
Total params: 256,571
Trainable params: 256,571
Non-trainable params: 0

```

#Training the Model

```
history = model.fit(x = x_train, y = y_train, epochs = 5, validation_split = 0.05)
```

```
Epoch 1/5
1014/1014 [=====] - 12s 12ms/step - loss: 0.2029 - accuracy: 0.9373 - val_loss: 0.1583 - val_accuracy: 0.9473
Epoch 2/5
1014/1014 [=====] - 11s 11ms/step - loss: 0.1271 - accuracy: 0.9624 - val_loss: 0.1521 - val_accuracy: 0.9467
Epoch 3/5
1014/1014 [=====] - 12s 11ms/step - loss: 0.0906 - accuracy: 0.9754 - val_loss: 0.1479 - val_accuracy: 0.9537
Epoch 4/5
1014/1014 [=====] - 12s 12ms/step - loss: 0.0714 - accuracy: 0.9801 - val_loss: 0.1644 - val_accuracy: 0.9525
Epoch 5/5
1014/1014 [=====] - 12s 12ms/step - loss: 0.0588 - accuracy: 0.9839 - val_loss: 0.1541 - val_accuracy: 0.9584
```

#Evaluation

```
evaluate = model.evaluate(x_test, y_test)
```

```
57/57 [=====] - 0s 3ms/step - loss: 0.1112 - accuracy: 0.9688
```

```
print("Test Acuracy is : {:.2f} %".format(evaluate[1]*100))
print("Test Loss is : {:.4f}".format(evaluate[0]))
```

```
Test Acuracy is : 96.88 %
Test Loss is : 0.1112
```

```
predictions = model.predict(x_test)
```

```
57/57 [=====] - 0s 2ms/step
```

```
predict = []
for i in predictions:
    predict.append(np.argmax(i))
```

Confusion Matrix

```
from sklearn import metrics
cm = metrics.confusion_matrix(predict,y_test)
acc = metrics.accuracy_score(predict,y_test)

print("The Confusion matrix is: \n",cm)
```

The Confusion matrix is:

```
[[548  22   1]
 [   6 572   8]
 [   0  19 621]]
```

#Accuracy

```
print(acc*100)
```

96.88369504730106

#Classification Report

```
from sklearn import metrics
print(metrics.classification_report(y_test, predict))
```

	precision	recall	f1-score	support
0	0.96	0.99	0.97	554
1	0.98	0.93	0.95	613
2	0.97	0.99	0.98	630
accuracy			0.97	1797
macro avg	0.97	0.97	0.97	1797
weighted avg	0.97	0.97	0.97	1797

#LSTM

```
from tensorflow.keras.layers import Embedding, LSTM, Dense
# ARCHITECTURE
EMBED_DIM = 32
LSTM_OUT = 64

model = Sequential()
model.add(Embedding(num_words, EMBED_DIM, input_length = pad_length))
model.add(LSTM(LSTM_OUT))
model.add(Dense(3, activation='softmax'))
model.compile(optimizer = 'adam', loss = 'sparse_categorical_crossentropy',
metrics = ['accuracy'])
```

```
print(model.summary())
```

Model: "sequential_4"

Layer (type)	Output Shape	Param #
embedding_4 (Embedding)	(None, 24, 32)	256000
lstm_1 (LSTM)	(None, 64)	24832
dense_5 (Dense)	(None, 3)	195

```
=====  
Total params: 281,027  
Trainable params: 281,027  
Non-trainable params: 0
```

None

#Training Model

```
history = model.fit(x = x_train, y = y_train, epochs = 10, validation_split = 0.05)
```

Epoch 1/10

```
1014/1014 [=====] - 26s 24ms/step - loss: 0.4328 - accuracy: 0.8246 - val_loss: 0.2531 - val_accuracy: 0.9127
```

Epoch 2/10

```
1014/1014 [=====] - 22s 22ms/step - loss: 0.1874 - accuracy: 0.9405 - val_loss: 0.1876 - val_accuracy: 0.9367
```

Epoch 3/10

```
1014/1014 [=====] - 23s 23ms/step - loss: 0.1426 - accuracy: 0.9554 - val_loss: 0.1991 - val_accuracy: 0.9391
```

Epoch 4/10

```
1014/1014 [=====] - 24s 23ms/step - loss: 0.1188 - accuracy: 0.9626 - val_loss: 0.1991 - val_accuracy: 0.9361
```

Epoch 5/10

```
1014/1014 [=====] - 23s 23ms/step - loss: 0.0999 - accuracy: 0.9688 - val_loss: 0.2142 - val_accuracy: 0.9408
```

Epoch 6/10

```
1014/1014 [=====] - 23s 22ms/step - loss: 0.0835 - accuracy: 0.9730 - val_loss: 0.1820 - val_accuracy: 0.9508
```

Epoch 7/10

```
1014/1014 [=====] - 24s 23ms/step - loss: 0.0735 - accuracy: 0.9767 - val_loss: 0.2202 - val_accuracy: 0.9473
```

Epoch 8/10

```
1014/1014 [=====] - 24s 23ms/step - loss: 0.0670 - accuracy: 0.9783 - val_loss: 0.2666 - val_accuracy: 0.9461
```

Epoch 9/10

```
1014/1014 [=====] - 22s 22ms/step - loss: 0.0577 - accuracy: 0.9815 - val_loss: 0.1988 - val_accuracy: 0.9490
```

Epoch 10/10

```
1014/1014 [=====] - 24s 23ms/step - loss: 0.0537 - accuracy: 0.9823 - val_loss: 0.2313 - val_accuracy: 0.9490
```

#Evaluation

```
evaluate = model.evaluate(x_test,y_test)
```

```
57/57 [=====] - 0s 8ms/step - loss: 0.1719 - accuracy: 0.9594
```

```
print("Test Acuracy is : {:.2f} %".format(evaluate[1]*100))
print("Test Loss is : {:.4f}".format(evaluate[0]))
```

```
Test Acuracy is : 95.94 %
Test Loss is : 0.1719
```

```
predictions = model.predict(x_test)
```

```
57/57 [=====] - 1s 8ms/step
```

```
predict = []
for i in predictions:
    predict.append(np.argmax(i))
```

#Confusion Matrix

```
from sklearn import metrics
cm = metrics.confusion_matrix(predict,y_test)
acc = metrics.accuracy_score(predict,y_test)
```

```
print("The Confusion matrix is: \n",cm)
```

The Confusion matrix is:

```
[[552  42   1]
 [  0 543   0]
 [  2  28 629]]
```

#Accuracy

```
print(acc*100)
```

```
95.93767390094602
```

#Classification Report

```
from sklearn import metrics
print(metrics.classification_report(y_test, predict))
```

	precision	recall	f1-score	support
0	0.93	1.00	0.96	554

1	1.00	0.89	0.94	613
2	0.95	1.00	0.98	630
accuracy			0.96	1797
macro avg	0.96	0.96	0.96	1797
weighted avg	0.96	0.96	0.96	1797

Experiment - 7

Auto encoders using MNIST data

Importing Necessary libraries

```
import numpy as np
import tensorflow as tf
import matplotlib.pyplot as plt

from tensorflow.keras import layers
from tensorflow.keras.datasets import mnist
from tensorflow.keras.models import Model
```

Pre-processing

```
def preprocess(array):
    array = array.astype("float32") / 255.0
    array = np.reshape(array, (len(array), 28, 28, 1))
    return array
```

Adding noise to the original images

```
def noise(array):
    noise_factor = 0.4 #amount of noise to add
    noisy_array = array + noise_factor * np.random.normal(
        loc=0.0, scale=1.0, size=array.shape
    )

    return np.clip(noisy_array, 0.0, 1.0)
```

Visualizing the images

```
def display(array1, array2):
    n = 10

    indices = np.random.randint(len(array1), size=n)
    images1 = array1[indices, :]
    images2 = array2[indices, :]

    plt.figure(figsize=(20, 4))
    for i, (image1, image2) in enumerate(zip(images1, images2)):
        ax = plt.subplot(2, n, i + 1)
        plt.imshow(image1.reshape(28, 28))
        plt.gray()
        ax.get_xaxis().set_visible(False)
```

```

ax.get_yaxis().set_visible(False)

ax = plt.subplot(2, n, i + 1 + n)
plt.imshow(image2.reshape(28, 28))
plt.gray()
ax.get_xaxis().set_visible(False)
ax.get_yaxis().set_visible(False)

plt.show()

```

Preparing the data

```
(train_data, _), (test_data, _) = mnist.load_data()
```

Normalize and reshape the data

```
train_data = preprocess(train_data)
```

```
test_data = preprocess(test_data)
```

Create a copy of the data with added noise

```
noisy_train_data = noise(train_data)
```


```
noisy_test_data = noise(test_data)
```

Display the train data and a version of it with added noise

```
display(train_data, noisy_train_data)
```

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz>

11490434/11490434 ————— 2s 0us/step



Building the Autoencoder

```
input = layers.Input(shape=(28, 28, 1))
```

Encoder

```
x = layers.Conv2D(32, (3, 3), activation="relu", padding="same")(input)
```

```
x = layers.MaxPooling2D((2, 2), padding="same")(x)
```

```
x = layers.Conv2D(32, (3, 3), activation="relu", padding="same")(x)
```

```
x = layers.MaxPooling2D((2, 2), padding="same")(x)
```

Decoder

```
x = layers.Conv2DTranspose(32, (3, 3), strides=2, activation="relu",
```

```
padding="same")(x)
x = layers.Conv2DTranspose(32, (3, 3), strides=2, activation="relu",
padding="same")(x)
x = layers.Conv2D(1, (3, 3), activation="sigmoid", padding="same")(x)
```

```
# Autoencoder
```

```
autoencoder = Model(input, x)
autoencoder.compile(optimizer="adam", loss="binary_crossentropy")
autoencoder.summary()
```

```
Model: "functional"
```

Layer (type) Param #	Output Shape
input_layer (InputLayer) 0	(None, 28, 28, 1)
conv2d (Conv2D) 320	(None, 28, 28, 32)
max_pooling2d (MaxPooling2D) 0	(None, 14, 14, 32)
conv2d_1 (Conv2D) 9,248	(None, 14, 14, 32)
max_pooling2d_1 (MaxPooling2D) 0	(None, 7, 7, 32)
conv2d_transpose (Conv2DTranspose) 9,248	(None, 14, 14, 32)
conv2d_transpose_1 (Conv2DTranspose) 9,248	(None, 28, 28, 32)
conv2d_2 (Conv2D) 289	(None, 28, 28, 1)

Total params: 28,353 (110.75 KB)

Trainable params: 28,353 (110.75 KB)

Non-trainable params: 0 (0.00 B)

Training the model

```
autoencoder.fit(  
    x=train_data,  
    y=train_data,  
    epochs=50,  
    batch_size=128,  
    shuffle=True,  
    validation_data=(test_data, test_data),  
)
```

Epoch 1/50			
469/469	_____	10s 12ms/step	- loss: 0.2439 - val_loss: 0.0743
Epoch 2/50			
469/469	_____	3s 6ms/step	- loss: 0.0736 - val_loss: 0.0701
Epoch 3/50			
469/469	_____	3s 6ms/step	- loss: 0.0704 - val_loss: 0.0685
Epoch 4/50			
469/469	_____	3s 6ms/step	- loss: 0.0689 - val_loss: 0.0676
Epoch 5/50			
469/469	_____	3s 6ms/step	- loss: 0.0679 - val_loss: 0.0669
Epoch 6/50			
469/469	_____	3s 6ms/step	- loss: 0.0672 - val_loss: 0.0663
Epoch 7/50			
469/469	_____	5s 6ms/step	- loss: 0.0667 - val_loss: 0.0659
Epoch 8/50			
469/469	_____	5s 6ms/step	- loss: 0.0663 - val_loss: 0.0656
Epoch 9/50			
469/469	_____	3s 6ms/step	- loss: 0.0659 - val_loss: 0.0652
Epoch 10/50			
469/469	_____	5s 6ms/step	- loss: 0.0656 - val_loss: 0.0649
Epoch 11/50			
469/469	_____	5s 6ms/step	- loss: 0.0654 - val_loss: 0.0647
Epoch 12/50			
469/469	_____	5s 6ms/step	- loss: 0.0650 - val_loss: 0.0644
Epoch 13/50			
469/469	_____	5s 6ms/step	- loss: 0.0648 - val_loss: 0.0644
Epoch 14/50			
469/469	_____	5s 6ms/step	- loss: 0.0645 - val_loss: 0.0641
Epoch 15/50			
469/469	_____	5s 7ms/step	- loss: 0.0643 - val_loss: 0.0639
Epoch 16/50			
469/469	_____	5s 6ms/step	- loss: 0.0642 - val_loss: 0.0637
Epoch 17/50			

469/469	_____	3s	6ms/step	-	loss: 0.0641	-	val_loss: 0.0636
Epoch 18/50							
469/469	_____	5s	6ms/step	-	loss: 0.0640	-	val_loss: 0.0634
Epoch 19/50							
469/469	_____	5s	6ms/step	-	loss: 0.0638	-	val_loss: 0.0634
Epoch 20/50							
469/469	_____	3s	6ms/step	-	loss: 0.0637	-	val_loss: 0.0633
Epoch 21/50							
469/469	_____	3s	7ms/step	-	loss: 0.0636	-	val_loss: 0.0632
Epoch 22/50							
469/469	_____	5s	6ms/step	-	loss: 0.0634	-	val_loss: 0.0631
Epoch 23/50							
469/469	_____	5s	6ms/step	-	loss: 0.0635	-	val_loss: 0.0630
Epoch 24/50							
469/469	_____	3s	7ms/step	-	loss: 0.0634	-	val_loss: 0.0629
Epoch 25/50							
469/469	_____	3s	6ms/step	-	loss: 0.0634	-	val_loss: 0.0629
Epoch 26/50							
469/469	_____	3s	6ms/step	-	loss: 0.0633	-	val_loss: 0.0629
Epoch 27/50							
469/469	_____	5s	7ms/step	-	loss: 0.0632	-	val_loss: 0.0628
Epoch 28/50							
469/469	_____	3s	6ms/step	-	loss: 0.0630	-	val_loss: 0.0627
Epoch 29/50							
469/469	_____	3s	6ms/step	-	loss: 0.0630	-	val_loss: 0.0626
Epoch 30/50							
469/469	_____	3s	6ms/step	-	loss: 0.0630	-	val_loss: 0.0626
Epoch 31/50							
469/469	_____	6s	7ms/step	-	loss: 0.0631	-	val_loss: 0.0626
Epoch 32/50							
469/469	_____	5s	6ms/step	-	loss: 0.0628	-	val_loss: 0.0626
Epoch 33/50							
469/469	_____	3s	6ms/step	-	loss: 0.0628	-	val_loss: 0.0625
Epoch 34/50							
469/469	_____	6s	7ms/step	-	loss: 0.0628	-	val_loss: 0.0625
Epoch 35/50							
469/469	_____	5s	6ms/step	-	loss: 0.0628	-	val_loss: 0.0625
Epoch 36/50							
469/469	_____	3s	6ms/step	-	loss: 0.0629	-	val_loss: 0.0624
Epoch 37/50							
469/469	_____	6s	7ms/step	-	loss: 0.0628	-	val_loss: 0.0625
Epoch 38/50							
469/469	_____	3s	6ms/step	-	loss: 0.0628	-	val_loss: 0.0625
Epoch 39/50							
469/469	_____	3s	6ms/step	-	loss: 0.0627	-	val_loss: 0.0624
Epoch 40/50							
469/469	_____	5s	7ms/step	-	loss: 0.0626	-	val_loss: 0.0623
Epoch 41/50							
469/469	_____	3s	6ms/step	-	loss: 0.0627	-	val_loss: 0.0622
Epoch 42/50							


```
469/469 _____ 3s 6ms/step - loss: 0.0626 - val_loss: 0.0622
Epoch 43/50
469/469 _____ 3s 6ms/step - loss: 0.0625 - val_loss: 0.0622
Epoch 44/50
469/469 _____ 3s 6ms/step - loss: 0.0626 - val_loss: 0.0622
Epoch 45/50
469/469 _____ 5s 6ms/step - loss: 0.0625 - val_loss: 0.0621
Epoch 46/50
469/469 _____ 3s 6ms/step - loss: 0.0625 - val_loss: 0.0622
Epoch 47/50
469/469 _____ 5s 6ms/step - loss: 0.0625 - val_loss: 0.0622
Epoch 48/50
469/469 _____ 5s 6ms/step - loss: 0.0625 - val_loss: 0.0621
Epoch 49/50
469/469 _____ 5s 6ms/step - loss: 0.0624 - val_loss: 0.0621
Epoch 50/50
469/469 _____ 3s 7ms/step - loss: 0.0624 - val_loss: 0.0620
```

<keras.src.callbacks.history.History at 0x79e5e6633490>

Prediction

```
predictions = autoencoder.predict(test_data)
display(test_data, predictions)
```

```
313/313 _____ 1s 2ms/step
```



Experiment - 8

Variational Autoencoders

Imports

For working with and visualizing the data

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

For training the VAE

```
import tensorflow as tf
```

For creating interactive widgets

```
import ipywidgets as widgets
from IPython.display import display
```

2024-05-22 06:28:22.848294: I

tensorflow/core/platform/cpu_feature_guard.cc:182] This TensorFlow binary is optimized to use available CPU instructions in performance-critical operations.

To enable the following instructions: AVX2 FMA, in other operations, rebuild TensorFlow with the appropriate compiler flags.

Load the data from a .csv file

```
pixel_data = pd.read_csv('/content/age_gender.csv')['pixels']
```

Shuffle the data

```
pixel_data = pixel_data.sample(frac=1.0, random_state=1)
```

Convert the data into a NumPy array

```
pixel_data = pixel_data.apply(lambda x: np.array(x.split(" "), dtype=np.int))
```

```
pixel_data = np.stack(np.array(pixel_data), axis=0)
```

Rescale pixel values to be between 0 and 1

```
pixel_data = pixel_data * (1./255)
```

/tmp/ipykernel_92174/3099532490.py:4: DeprecationWarning: `np.int` is a deprecated alias for the builtin `int`. To silence this warning, use `int` by itself. Doing this will not modify any behavior and is safe. When replacing `np.int`, you may wish to use e.g. `np.int64` or `np.int32` to specify the precision. If you wish to review your current use, check the release note link for additional information.

Deprecated in NumPy 1.20; for more details and guidance:

<https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations>

```
pixel_data = pixel_data.apply(lambda x: np.array(x.split(" "),
dtype=np.int))
```

```
# The data is now a NumPy array of 23705 images.
```

```
# we are working with 48x48x1 images)
```

```
pixel_data.shape
```

```
(23705, 2304)
```

Building the VAE

```
class Sampling(tf.keras.layers.Layer):
```

```
    def call(self, inputs):
```

```
        z_mean, z_log_var = inputs # Unpack the inputs into mean and log-
variance
```

```
        batch = tf.shape(z_mean)[0] # Get the batch size
```

```
        dim = tf.shape(z_mean)[1] # Get the dimensionality of the latent
space
```

```
        epsilon = tf.keras.backend.random_normal(shape=(batch, dim)) #
Sample from standard normal distribution
```

```
        return epsilon * tf.exp(z_log_var * 0.5) + z_mean # Apply the
reparameterization trick
```

```
def build_vae(num_pixels, num_latent_vars=3):
```

```
    # Encoder
```

```
    encoder_inputs = tf.keras.Input(shape=(num_pixels,)) # Input layer for
the encoder
```

```
    x = tf.keras.layers.Dense(512, activation='relu')(encoder_inputs) #
First dense layer with 512 units and ReLU activation
```

```
    x = tf.keras.layers.Dense(128, activation='relu')(x) # Second dense
layer with 128 units and ReLU activation
```

```
    x = tf.keras.layers.Dense(32, activation='relu')(x) # Third dense layer
with 32 units and ReLU activation
```

```
    z_mean = tf.keras.layers.Dense(num_latent_vars)(x) # Dense layer for the
mean of the latent variables
```

```
    z_log_var = tf.keras.layers.Dense(num_latent_vars)(z_mean) # Dense layer
for the log-variance of the latent variables
```

```
    z = Sampling()([z_mean, z_log_var]) # Sampling layer to sample the
latent variables using the reparameterization trick
```

```
    encoder = tf.keras.Model(inputs=encoder_inputs, outputs=z) # Define the
encoder model
```

```
    # Decoder
```

```
    decoder_inputs = tf.keras.Input(shape=(num_latent_vars,)) # Input layer
for the decoder
```

```
    x = tf.keras.layers.Dense(32, activation='relu')(decoder_inputs) # First
dense layer with 32 units and ReLU activation
```

```
    x = tf.keras.layers.Dense(128, activation='relu')(x) # Second dense
```

```

Layer with 128 units and ReLU activation
    x = tf.keras.layers.Dense(512, activation='relu')(x) # Third dense layer
with 512 units and ReLU activation
    reconstruction = tf.keras.layers.Dense(num_pixels,
activation='linear')(x) # Output dense layer with 'num_pixels' units and
Linear activation

    decoder = tf.keras.Model(inputs=decoder_inputs, outputs=reconstruction)
# Define the decoder model

    # Full model
    model_inputs = encoder.input # Inputs of the full VAE model are the
inputs of the encoder
    model_outputs = decoder(encoder.output) # Outputs of the full VAE model
are the outputs of the decoder, given the encoder's output

    model = tf.keras.Model(inputs=model_inputs, outputs=model_outputs) #
Define the full VAE model

    # Compile model for training
    model.compile(
        optimizer='adam', # Adam optimizer
        loss='mse' # Mean Squared Error (MSE) Loss function
    )

    # Return all three models
    return encoder, decoder, model # Return the encoder, decoder, and full
VAE models

```

```

face_encoder, face_decoder, face_model = build_vae(num_pixels=2304,
num_latent_vars=3)

```

```

2024-05-22 06:28:31.067844: W
tensorflow/core/common_runtime/gpu/gpu_device.cc:1960] Cannot dlopen some GPU
libraries. Please make sure the missing libraries mentioned above are
installed properly if you would like to use GPU. Follow the guide at
https://www.tensorflow.org/install/gpu for how to download and setup the
required libraries for your platform.
Skipping registering GPU devices...

```

```

print(face_encoder.summary())

```

```

Model: "model"

```

Layer (type)	Output Shape	Param #	Connected to
=====			
=====			
input_1 (InputLayer)	[(None, 2304)]	0	[]

dense (Dense) ['input_1[0][0]']	(None, 512)	1180160
dense_1 (Dense) ['dense[0][0]']	(None, 128)	65664
dense_2 (Dense) ['dense_1[0][0]']	(None, 32)	4128
dense_3 (Dense) ['dense_2[0][0]']	(None, 3)	99
dense_4 (Dense) ['dense_3[0][0]']	(None, 3)	12
sampling (Sampling) ['dense_3[0][0]', 'dense_4[0][0]']	(None, 3)	0

```
=====
Total params: 1250063 (4.77 MB)
Trainable params: 1250063 (4.77 MB)
Non-trainable params: 0 (0.00 Byte)
```

None

```
print(face_decoder.summary())
```

Model: "model_1"

Layer (type)	Output Shape	Param #
input_2 (InputLayer)	[(None, 3)]	0
dense_5 (Dense)	(None, 32)	128
dense_6 (Dense)	(None, 128)	4224
dense_7 (Dense)	(None, 512)	66048
dense_8 (Dense)	(None, 2304)	1181952

```
=====
Total params: 1252352 (4.78 MB)
Trainable params: 1252352 (4.78 MB)
```

Non-trainable params: 0 (0.00 Byte)

None

Train the VAE

```
history = face_model.fit(
    pixel_data,
    pixel_data,
    validation_split=0.2,
    batch_size=32,
    epochs=100,
    callbacks=[
        tf.keras.callbacks.EarlyStopping(
            monitor='val_loss',
            patience=10,
            restore_best_weights=True
        )
    ]
)
```

Epoch 1/100

593/593 [=====] - 8s 12ms/step - loss: 0.0256 - val_loss: 0.0216

Epoch 2/100

593/593 [=====] - 5s 9ms/step - loss: 0.0216 - val_loss: 0.0213

Epoch 3/100

593/593 [=====] - 5s 9ms/step - loss: 0.0214 - val_loss: 0.0212

Epoch 4/100

593/593 [=====] - 5s 9ms/step - loss: 0.0212 - val_loss: 0.0212

Epoch 5/100

593/593 [=====] - 5s 9ms/step - loss: 0.0212 - val_loss: 0.0210

Epoch 6/100

593/593 [=====] - 5s 9ms/step - loss: 0.0211 - val_loss: 0.0213

Epoch 7/100

593/593 [=====] - 5s 9ms/step - loss: 0.0210 - val_loss: 0.0209

Epoch 8/100

593/593 [=====] - 5s 9ms/step - loss: 0.0209 - val_loss: 0.0209

Epoch 9/100

593/593 [=====] - 5s 9ms/step - loss: 0.0209 - val_loss: 0.0210

Epoch 10/100

593/593 [=====] - 5s 9ms/step - loss: 0.0208 -


```
val_loss: 0.0206
Epoch 11/100
593/593 [=====] - 5s 9ms/step - loss: 0.0206 -
val_loss: 0.0206
Epoch 12/100
593/593 [=====] - 5s 9ms/step - loss: 0.0206 -
val_loss: 0.0208
Epoch 13/100
593/593 [=====] - 6s 9ms/step - loss: 0.0206 -
val_loss: 0.0207
Epoch 14/100
593/593 [=====] - 5s 9ms/step - loss: 0.0206 -
val_loss: 0.0205
Epoch 15/100
593/593 [=====] - 5s 9ms/step - loss: 0.0207 -
val_loss: 0.0210
Epoch 16/100
593/593 [=====] - 5s 9ms/step - loss: 0.0207 -
val_loss: 0.0204
Epoch 17/100
593/593 [=====] - 5s 9ms/step - loss: 0.0208 -
val_loss: 0.0208
Epoch 18/100
593/593 [=====] - 5s 9ms/step - loss: 0.0206 -
val_loss: 0.0206
Epoch 19/100
593/593 [=====] - 5s 9ms/step - loss: 0.0205 -
val_loss: 0.0207
Epoch 20/100
593/593 [=====] - 5s 9ms/step - loss: 0.0207 -
val_loss: 0.0205
Epoch 21/100
593/593 [=====] - 5s 9ms/step - loss: 0.0207 -
val_loss: 0.0206
Epoch 22/100
593/593 [=====] - 5s 9ms/step - loss: 0.0207 -
val_loss: 0.0209
Epoch 23/100
593/593 [=====] - 5s 9ms/step - loss: 0.0207 -
val_loss: 0.0207
Epoch 24/100
593/593 [=====] - 5s 9ms/step - loss: 0.0206 -
val_loss: 0.0210
Epoch 25/100
593/593 [=====] - 5s 9ms/step - loss: 0.0205 -
val_loss: 0.0205
Epoch 26/100
593/593 [=====] - 5s 9ms/step - loss: 0.0205 -
val_loss: 0.0205
```

Image Reconstruction

i = 6

```
sample = np.array(pixel_data)[i].copy()
sample = sample.reshape(48, 48, 1)

reconstruction = face_model.predict(pixel_data)[i].copy()
reconstruction = reconstruction.reshape(48, 48, 1)

plt.figure(figsize=(10, 5))

plt.subplot(1, 2, 1)
plt.imshow(sample, cmap='gray')
plt.axis('off')
plt.title("Original Image")

plt.subplot(1, 2, 2)
plt.imshow(reconstruction, cmap='gray')
plt.axis('off')
plt.title("Reconstructed Image")

plt.show()
```

741/741 [=====] - 1s 2ms/step

Original Image



Reconstructed Image



Specify our own latent variable values

```
def generate_face_image(latent1, latent2, latent3):
    latent_vars = np.array([[latent1, latent2, latent3]])
    reconstruction = np.array(face_decoder(latent_vars))
```

```

reconstruction = reconstruction.reshape(48, 48, 1)
plt.figure()
plt.imshow(reconstruction, cmap='gray')
plt.axis('off')
plt.show()

# Let's get the min and max for each slider on the interactive widget
latent1_min = np.min(face_encoder(pixel_data).numpy()[ :, 0])
latent1_max = np.max(face_encoder(pixel_data).numpy()[ :, 0])

latent2_min = np.min(face_encoder(pixel_data).numpy()[ :, 1])
latent2_max = np.max(face_encoder(pixel_data).numpy()[ :, 1])

latent3_min = np.min(face_encoder(pixel_data).numpy()[ :, 2])
latent3_max = np.max(face_encoder(pixel_data).numpy()[ :, 2])

import tensorflow as tf
print(tf.__version__)

2.13.1

# Create the interactive widget
face_image_generator = widgets.interact(
    generate_face_image,
    latent1=(latent1_min, latent1_max),
    latent2=(latent2_min, latent2_max),
    latent3=(latent3_min, latent3_max),
)

# Display the widget
display(face_image_generator)

```



```
{"model_id":"94a81f02f9594f228f804939dcececfe","version_major":2,"version_min  
or":0}
```

```
<function __main__.generate_face_image(latent1, latent2, latent3)>
```

Experiment - 9

Implementing GAN Architecture on MINST dataset

Importing Necessary libraries

```
from tensorflow.keras.layers import (Dense,
                                      BatchNormalization,
                                      LeakyReLU,
                                      Reshape,
                                      Conv2DTranspose,
                                      Conv2D,
                                      Dropout,
                                      Flatten)

import tensorflow as tf
import matplotlib.pyplot as plt
```

Preparing the data

```
(train_images, train_labels), (_, _) = tf.keras.datasets.mnist.load_data()
train_images = train_images.reshape(train_images.shape[0], 28, 28,
1).astype('float32')
train_images = (train_images - 127.5) / 127.5

BUFFER_SIZE = 60000
BATCH_SIZE = 256

# Batch and shuffle the data
train_dataset =
tf.data.Dataset.from_tensor_slices(train_images).shuffle(BUFFER_SIZE).batch(B
ATCH_SIZE)
```

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz>
11490434/11490434 [=====] - 0s 0us/step

Dimension of the Noise

```
z_dim = 100
```

Functional Code Generator

```
nch = 200
g_input = Input(shape=[100])
```

```

H = Dense(nch*14*14, kernel_initializer='glorot_normal')(g_input)
H = BatchNormalization()(H)
H = Activation('relu')(H)
H = Reshape( [nch, 14, 14] )(H)
H = UpSampling2D(size=(2, 2))(H)
H = Convolution2D(int(nch/2), 3, 3, padding='same',
kernel_initializer='glorot_uniform')(H)
H = BatchNormalization()(H)
H = Activation('relu')(H)
H = Convolution2D(int(nch/4), 3, 3, padding='same',
kernel_initializer='glorot_uniform')(H)
H = BatchNormalization()(H)
H = Activation('relu')(H)
H = Convolution2D(1, 1, 1, padding='same',
kernel_initializer='glorot_uniform')(H)
g_V = Activation('sigmoid')(H)
generator = Model(g_input,g_V)
generator.compile(loss='binary_crossentropy', optimizer=Adam(lr=0.0002,
beta_1=0.5))
generator.summary()

```

Model: "model"

Layer (type)	Output Shape	Param #
=====		
input_4 (InputLayer)	[(None, 100)]	0
dense_2 (Dense)	(None, 39200)	3959200
batch_normalization_1 (Batch Normalization)	(None, 39200)	156800
activation_1 (Activation)	(None, 39200)	0
reshape_1 (Reshape)	(None, 200, 14, 14)	0
up_sampling2d_1 (UpSampling 2D)	(None, 400, 28, 14)	0
conv2d (Conv2D)	(None, 134, 10, 100)	12700
batch_normalization_2 (Batch Normalization)	(None, 134, 10, 100)	400
activation_2 (Activation)	(None, 134, 10, 100)	0
conv2d_1 (Conv2D)	(None, 45, 4, 50)	45050
batch_normalization_3 (Batch Normalization)	(None, 45, 4, 50)	200

hNormalization)

activation_3 (Activation)	(None, 45, 4, 50)	0
conv2d_2 (Conv2D)	(None, 45, 4, 1)	51
activation_4 (Activation)	(None, 45, 4, 1)	0

```
=====
Total params: 4,174,401
Trainable params: 4,095,701
Non-trainable params: 78,700
```

```
/usr/local/lib/python3.10/dist-packages/keras/optimizers/legacy/adam.py:117:
UserWarning: The `lr` argument is deprecated, use `learning_rate` instead.
  super().__init__(name, **kwargs)
```

Building Generator

```
def generator_model():
    model = tf.keras.Sequential()
    model.add(Dense(7*7*256, use_bias=False, input_shape=(100,)))
    model.add(BatchNormalization())
    model.add(LeakyReLU())

    model.add(Reshape((7, 7, 256)))
    assert model.output_shape == (None, 7, 7, 256) # Note: None is the batch
size

    model.add(Conv2DTranspose(128, (5, 5), strides=(1, 1), padding='same',
use_bias=False))
    assert model.output_shape == (None, 7, 7, 128)
    model.add(BatchNormalization())
    model.add(LeakyReLU())

    model.add(Conv2DTranspose(64, (5, 5), strides=(2, 2), padding='same',
use_bias=False))
    assert model.output_shape == (None, 14, 14, 64)
    model.add(BatchNormalization())
    model.add(LeakyReLU())

    model.add(Conv2DTranspose(1, (5, 5), strides=(2, 2), padding='same',
use_bias=False, activation='tanh'))
    assert model.output_shape == (None, 28, 28, 1)

    print(model.summary())

    return model
```

```
generator = generator_model()
```

```
Model: "sequential"
```

Layer (type)	Output Shape	Param #
=====		
dense (Dense)	(None, 12544)	1254400
batch_normalization (Batch Normalization)	(None, 12544)	50176
leaky_re_lu (LeakyReLU)	(None, 12544)	0
reshape (Reshape)	(None, 7, 7, 256)	0
conv2d_transpose (Conv2DTranspose)	(None, 7, 7, 128)	819200
batch_normalization_1 (Batch Normalization)	(None, 7, 7, 128)	512
leaky_re_lu_1 (LeakyReLU)	(None, 7, 7, 128)	0
conv2d_transpose_1 (Conv2DTranspose)	(None, 14, 14, 64)	204800
batch_normalization_2 (Batch Normalization)	(None, 14, 14, 64)	256
leaky_re_lu_2 (LeakyReLU)	(None, 14, 14, 64)	0
conv2d_transpose_2 (Conv2DTranspose)	(None, 28, 28, 1)	1600


```
=====
Total params: 2,330,944
Trainable params: 2,305,472
Non-trainable params: 25,472
```

```
None
```

Generate a Sample Image

```
# Create a random noise and generate a sample
noise = tf.random.normal([1, 100])
generated_image = generator(noise, training=False)
# Visualize the generated sample
plt.imshow(generated_image[0, :, :, 0], cmap='gray')
```


<matplotlib.image.AxesImage at 0x7fa67e014b50>



#Discriminative Model

```
def discriminator_model():
    model = tf.keras.Sequential()

    model.add(Conv2D(64, (5, 5), strides=(2, 2), padding='same',
input_shape=[28, 28, 1]))
    model.add(LeakyReLU())
    model.add(Dropout(0.3))

    model.add(Conv2D(128, (5, 5), strides=(2, 2), padding='same'))
    model.add(LeakyReLU())
    model.add(Dropout(0.3))

    model.add(Flatten())
    model.add(Dense(1))

    print(model.summary())

    return model

discriminator = discriminator_model()
```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 14, 14, 64)	1664
leaky_re_lu_4 (LeakyReLU)	(None, 14, 14, 64)	0
dropout (Dropout)	(None, 14, 14, 64)	0
conv2d_2 (Conv2D)	(None, 7, 7, 128)	204928
leaky_re_lu_5 (LeakyReLU)	(None, 7, 7, 128)	0
dropout_1 (Dropout)	(None, 7, 7, 128)	0
flatten (Flatten)	(None, 6272)	0
dense_1 (Dense)	(None, 1)	6273

=====
Total params: 212,865
Trainable params: 212,865
Non-trainable params: 0

None

Configure the Model

```
cross_entropy = tf.keras.losses.BinaryCrossentropy(from_logits=True)
```

```
def discriminator_loss(real_output, fake_output):  
    real_loss = cross_entropy(tf.ones_like(real_output), real_output)  
    fake_loss = cross_entropy(tf.zeros_like(fake_output), fake_output)  
    total_loss = real_loss + fake_loss  
    return total_loss
```

```
def generator_loss(fake_output):  
    return cross_entropy(tf.ones_like(fake_output), fake_output)
```

```
generator_optimizer = tf.keras.optimizers.Adam(1e-4)  
discriminator_optimizer = tf.keras.optimizers.Adam(1e-4)
```

#Training

```
import os  
checkpoint_dir = '/content/GAN/'  
checkpoint_prefix = os.path.join(checkpoint_dir, "ckpt")  
checkpoint = tf.train.Checkpoint(generator_optimizer=generator_optimizer,
```

```
discriminator_optimizer=discriminator_optimizer,
                        generator=generator,
                        discriminator=discriminator)
```

```
EPOCHS = 60
```

```
num_examples_to_generate = 16
noise_dim = 100
seed = tf.random.normal([num_examples_to_generate, noise_dim])
```

#Training Steps

```
@tf.function
def train_step(images):

    noise = tf.random.normal([BATCH_SIZE, noise_dim])

    with tf.GradientTape() as gen_tape, tf.GradientTape() as disc_tape:
        generated_images = generator(noise, training=True)

        real_output = discriminator(images, training=True)
        fake_output = discriminator(generated_images, training=True)

        gen_loss = generator_loss(fake_output)
        disc_loss = discriminator_loss(real_output, fake_output)

        gradients_of_generator = gen_tape.gradient(gen_loss,
                                                    generator.trainable_variables)
        gradients_of_discriminator = disc_tape.gradient(disc_loss,
                                                         discriminator.trainable_variables)

        generator_optimizer.apply_gradients(zip(gradients_of_generator,
                                                generator.trainable_variables))
        discriminator_optimizer.apply_gradients(zip(gradients_of_discriminator,
                                                    discriminator.trainable_variables))
```

#Image Generation Function

```
def generate_and_save_images(model, epoch, test_input):
    predictions = model(test_input, training=False)
    fig = plt.figure(figsize=(4,4))
    for i in range(predictions.shape[0]):
        plt.subplot(4, 4, i+1)
        plt.imshow(predictions[i, :, :, 0] * 127.5 + 127.5, cmap='gray')
        plt.axis('off')
```

```
plt.savefig('image_at_epoch_{:04d}.png'.format(epoch))
plt.show()
```

#Train GAN

```
import time
from IPython import display
def train(dataset, epochs):

    for epoch in range(epochs):
        start = time.time()

        for image_batch in dataset:
            train_step(image_batch)

        display.clear_output(wait=True)
        generate_and_save_images(generator,
                                epoch + 1,
                                seed)
    if (epoch + 1) % 5 == 0:
        checkpoint.save(file_prefix = checkpoint_prefix)

    print ('Time for epoch {} is {} sec'.format(epoch + 1, time.time()-
start))

    display.clear_output(wait=True)
    generate_and_save_images(generator,
                            epochs,
                            seed)
```

#Start Training

```
train(train_dataset, EPOCHS)
```



Generated Digits

```
checkpoint.restore(tf.train.latest_checkpoint(checkpoint_dir))
```

```
<tensorflow.python.checkpoint.checkpoint.CheckpointLoadStatus at  
0x7fa5c2d49450>
```

```
import PIL
```

```
def display_image(epoch_no):
```

```
    return PIL.Image.open('image_at_epoch_{:04d}.png'.format(epoch_no))
```

```
display_image(EPOCHS)
```



```
import glob # The glob module is used for Unix style pathname pattern
expansion.
import imageio # The library that provides an easy interface to read and
write a wide range of image data
```

```
anim_file = 'dcgan.gif'
```

```
with imageio.get_writer(anim_file, mode='I') as writer:
    filenames = glob.glob('image*.png')
    filenames = sorted(filenames)
    for filename in filenames:
        image = imageio.imread(filename)
        writer.append_data(image)
    # image = imageio.imread(filename)
    # writer.append_data(image)
```

```
display.Image(open('dcgan.gif', 'rb').read())
```